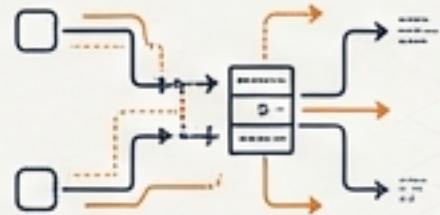


# Apache Parquet: The Anatomy of Efficiency

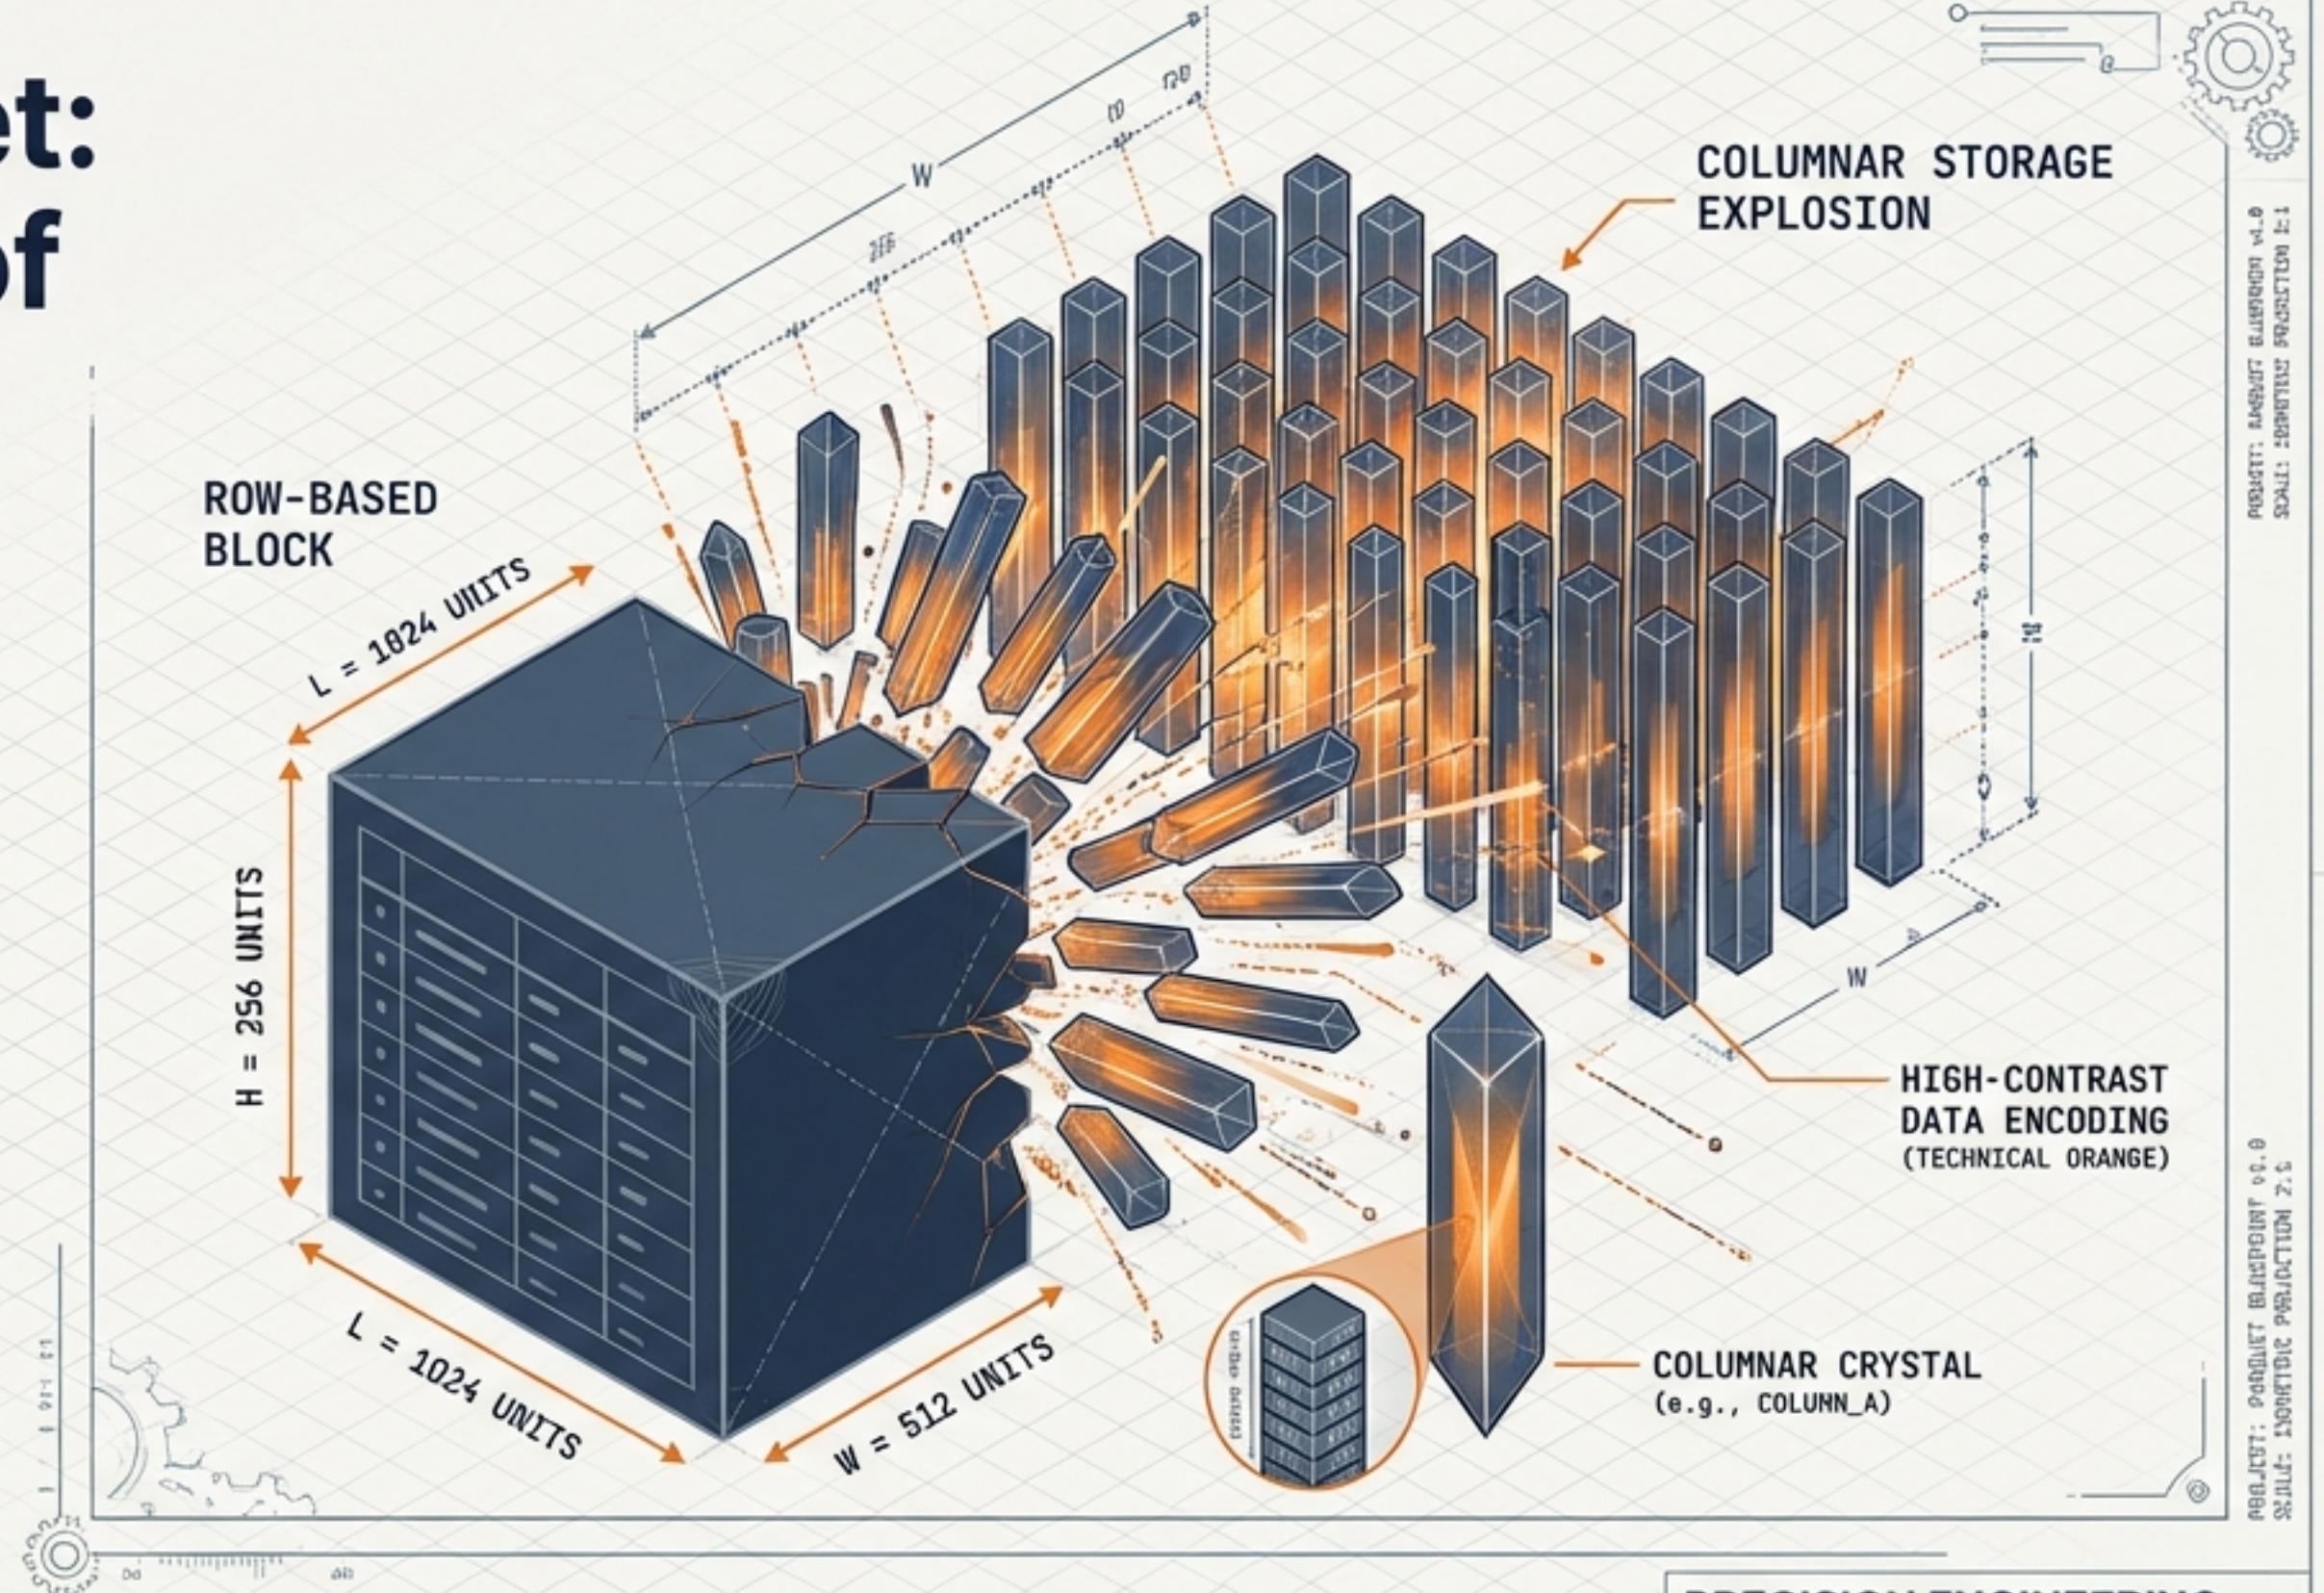
A Deep Dive into the Standard  
for Modern Data Analytics



Optimised for the  
Data Lakehouse

Engineered for Speed

The Industry Standard



PRECISION ENGINEERING  
BLUEPRINT

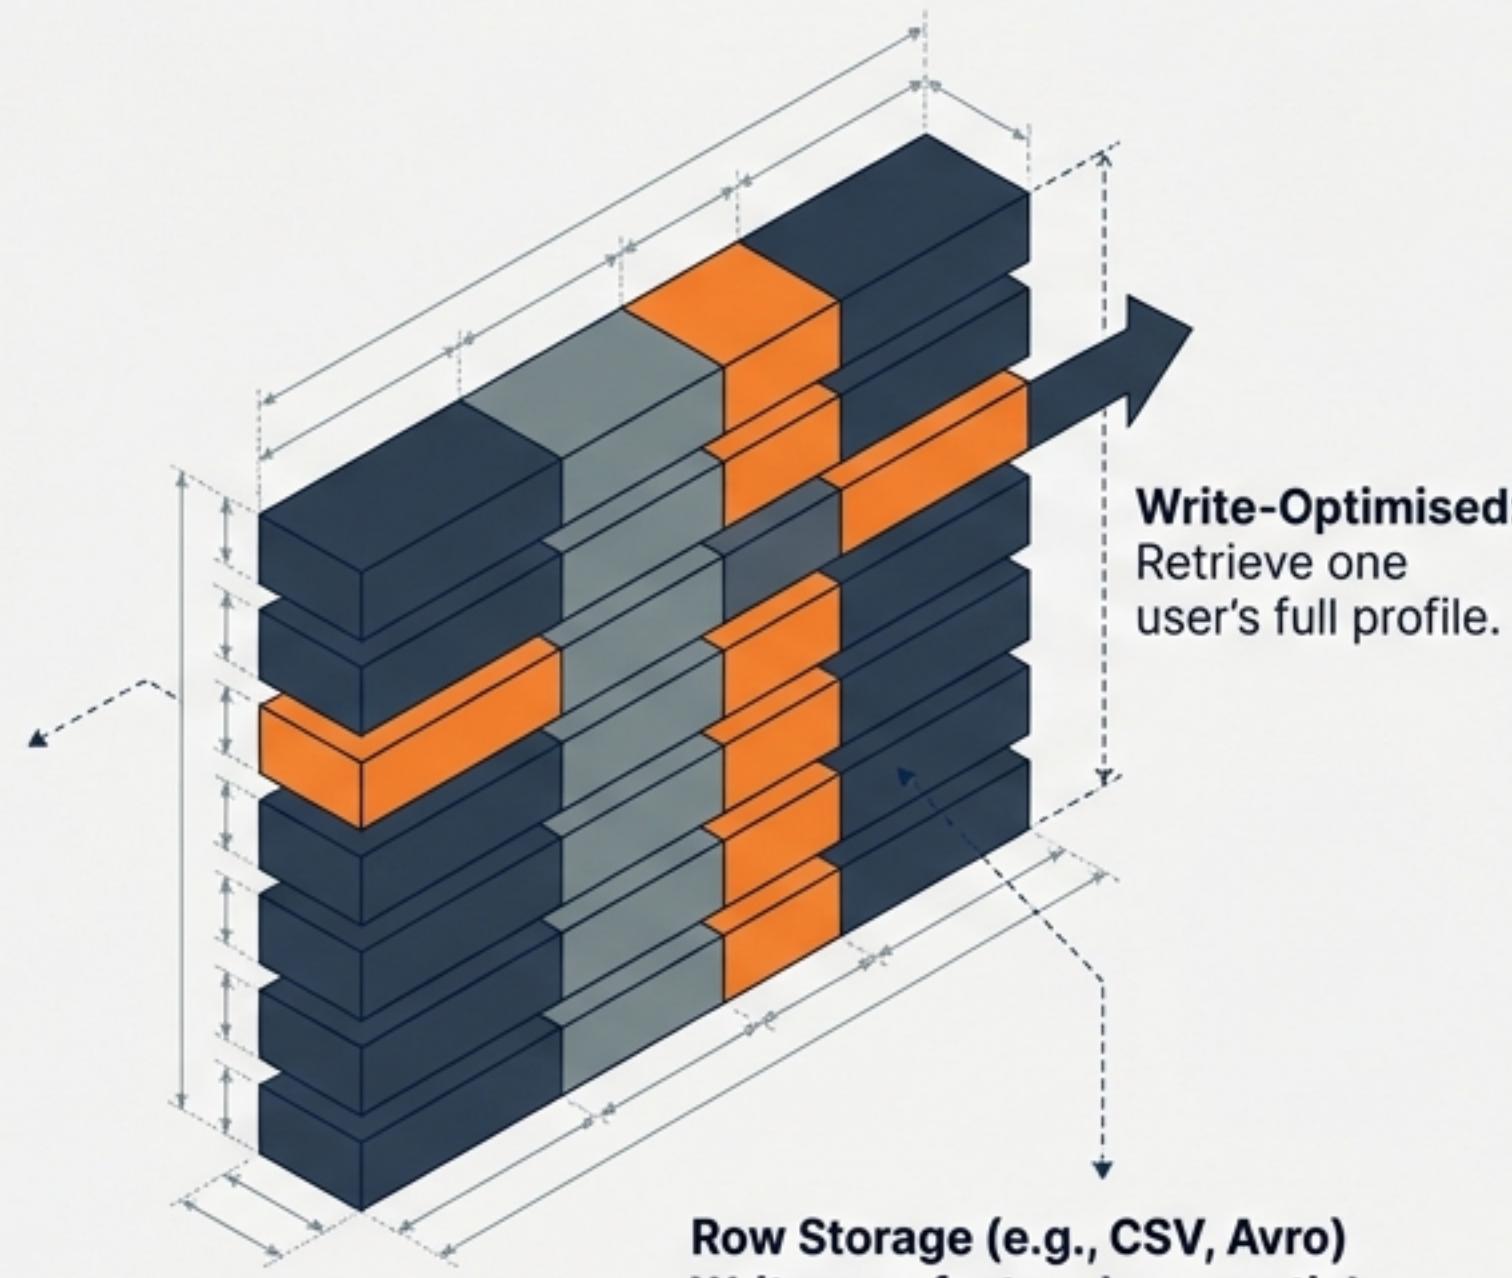
PROJECT: PARQUET BLUEPRINT v1.0  
SCALE: ISOMETRIC PROJECTION 2:1

Helvetica Now Display  
Helvetica Now Display

APPROVED FOR  
DATA ANALYTICS

# The Paradigm Shift: From Transactions to Analytics

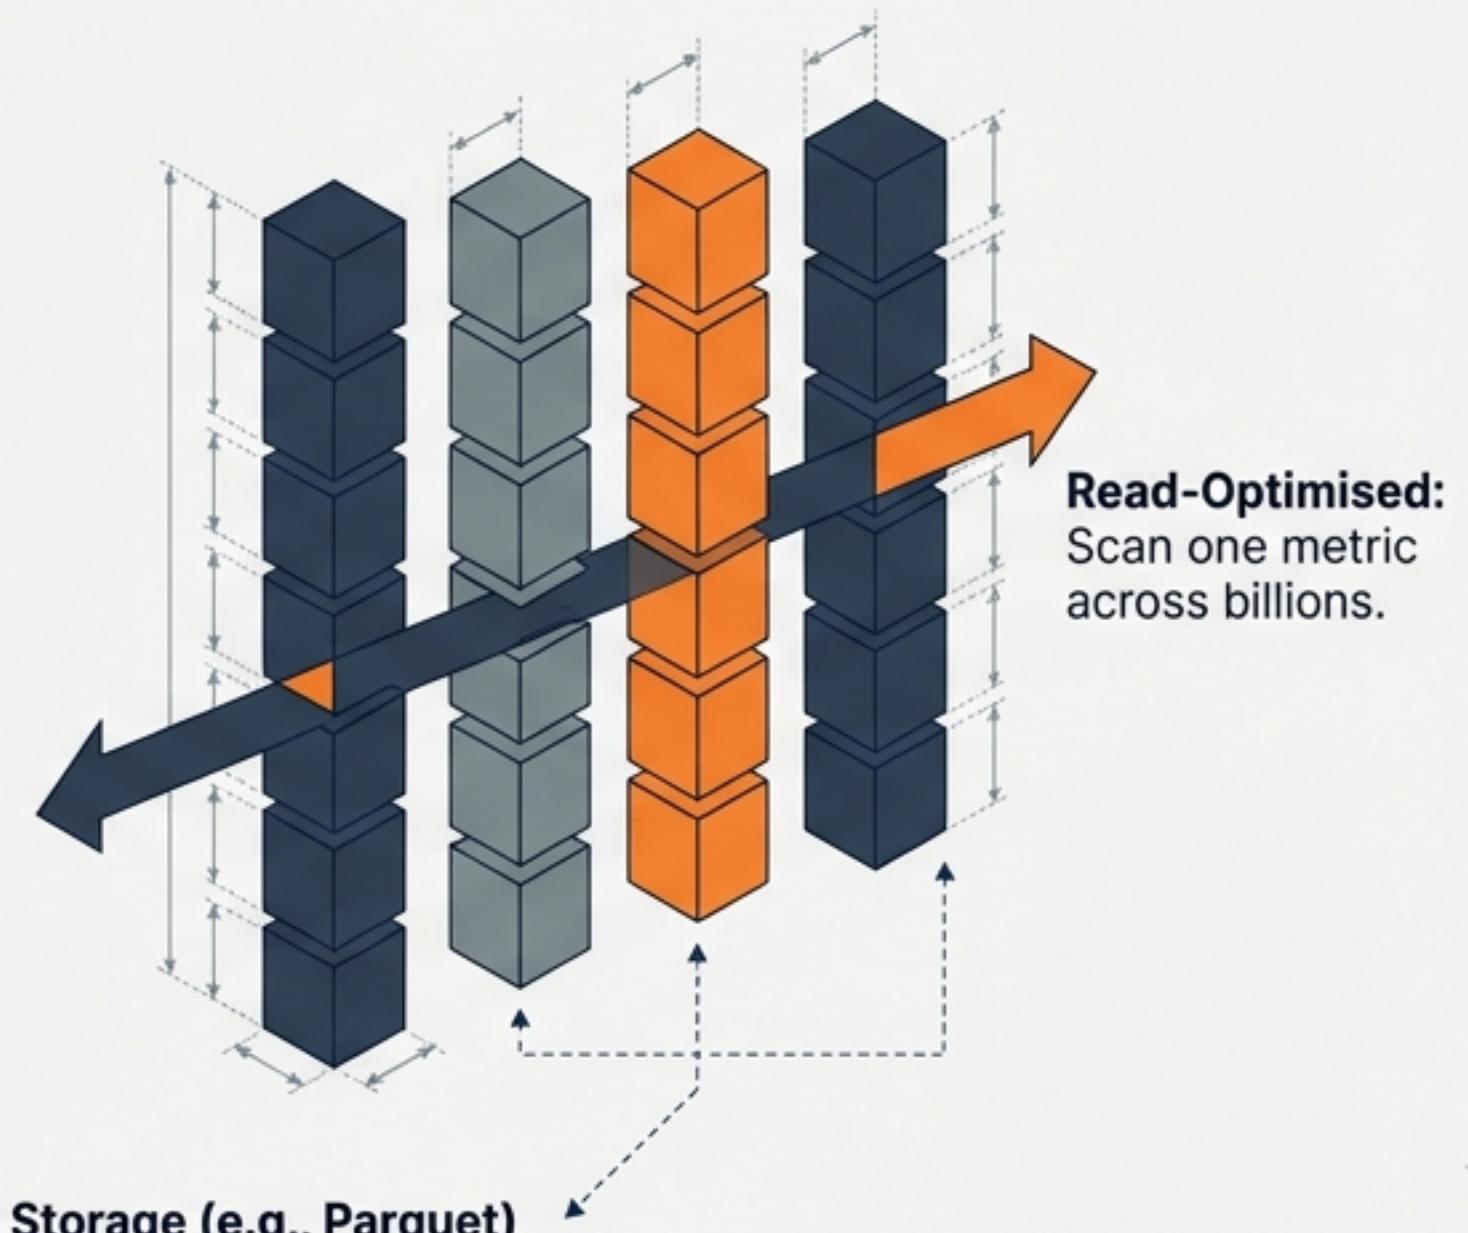
## Row-Oriented / OLTP



**Write-Optimised:**  
Retrieve one user's full profile.

**Row Storage (e.g., CSV, Avro)**  
**Writes are fast and sequential.**  
Best for transactional workloads where you need to retrieve one specific record's entire set of attributes.

## Column-Oriented / OLAP



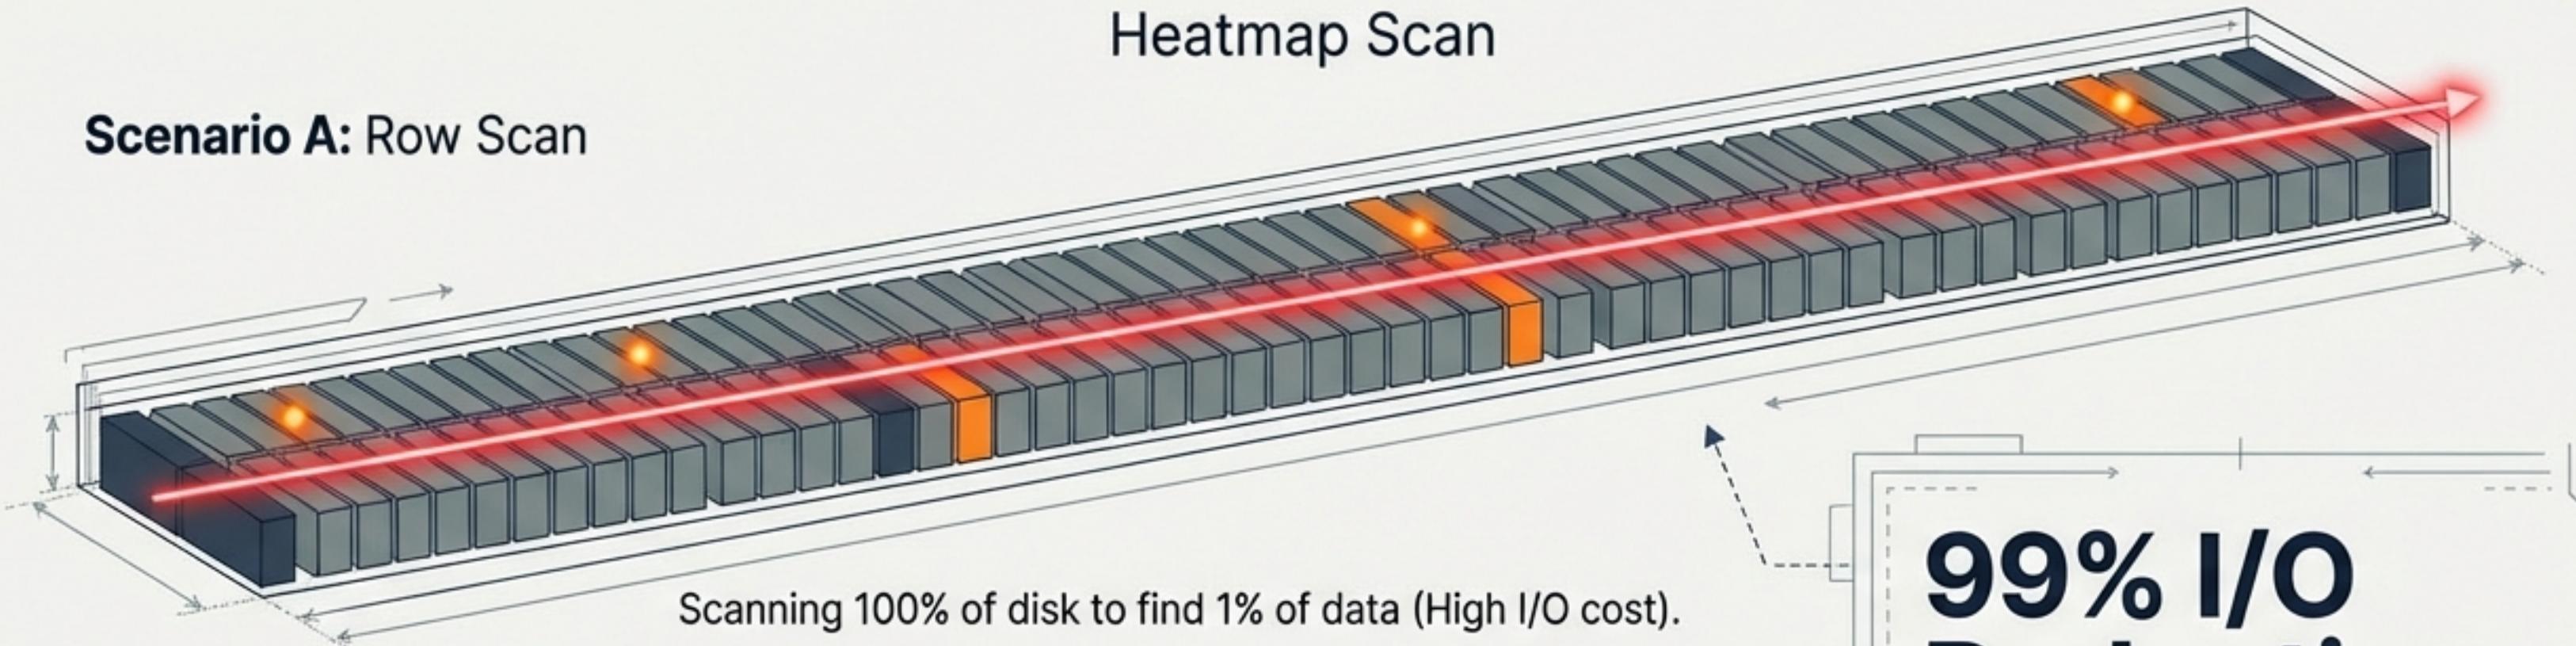
**Read-Optimised:**  
Scan one metric across billions.

**Column Storage (e.g., Parquet)**  
**Reads are selective.**  
Best for analytical workloads where you need to compute a metric (like Average Age) across millions of users without loading their Names or Addresses.

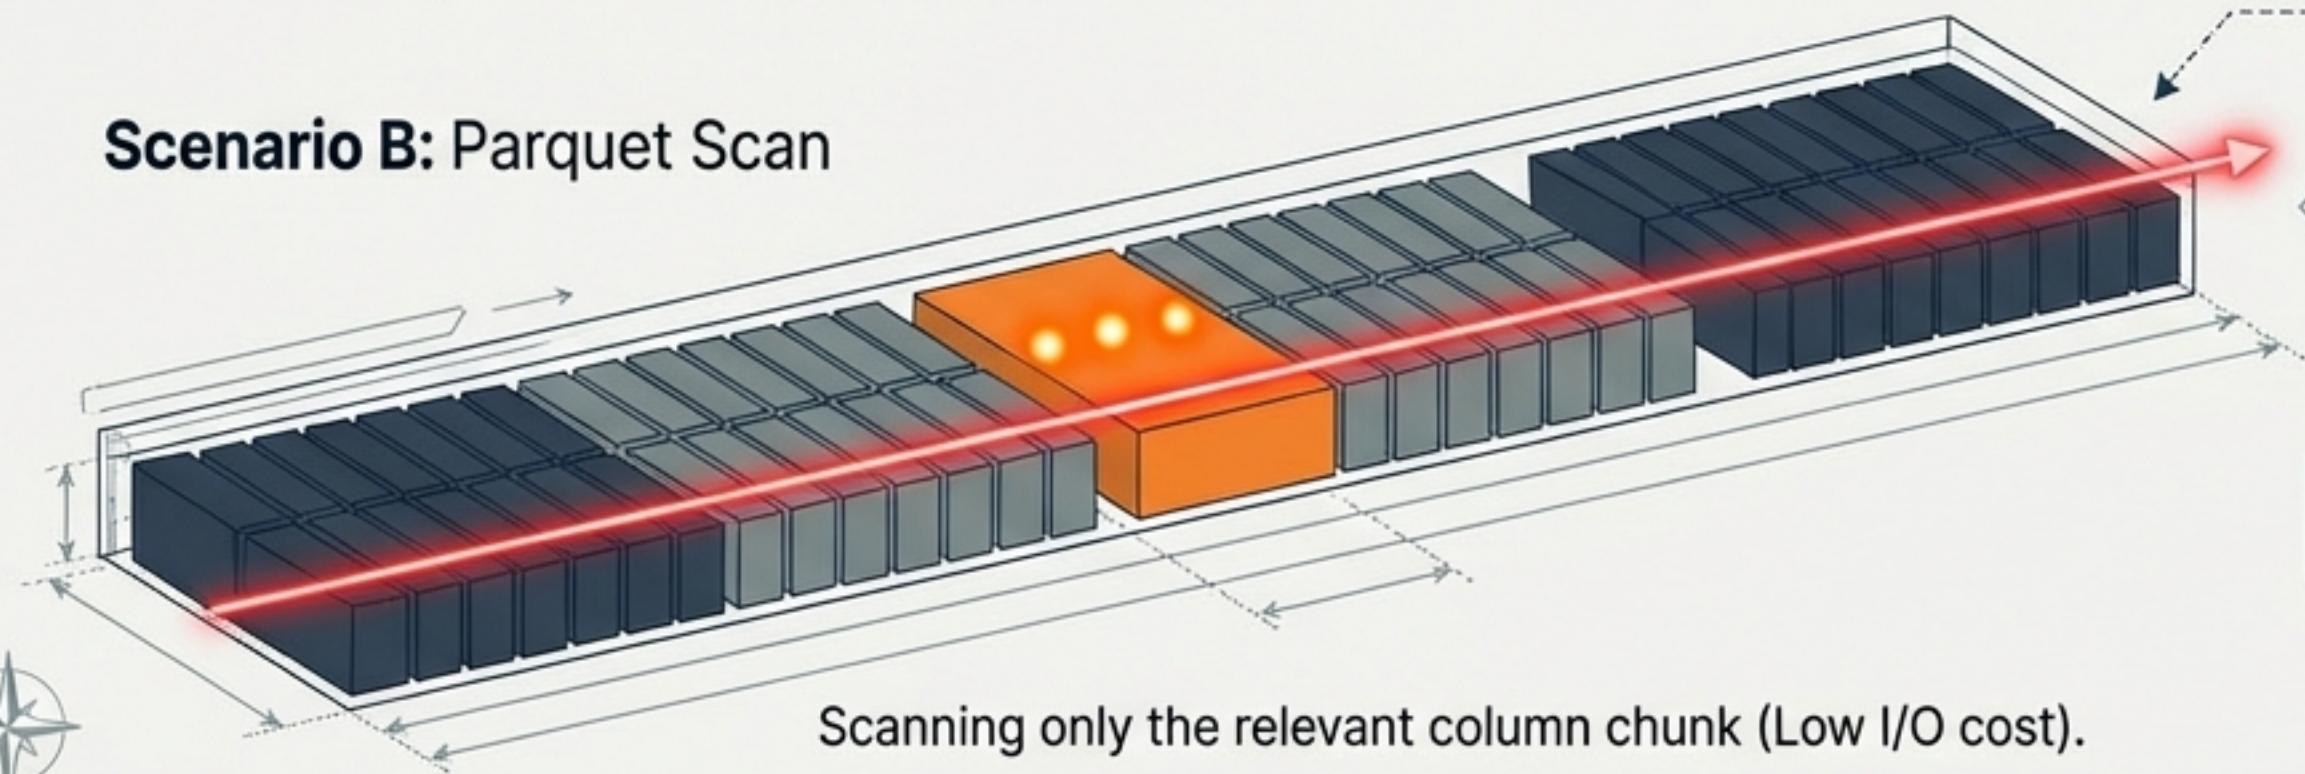
# The I/O Arbitrage

Heatmap Scan

**Scenario A: Row Scan**



**Scenario B: Parquet Scan**

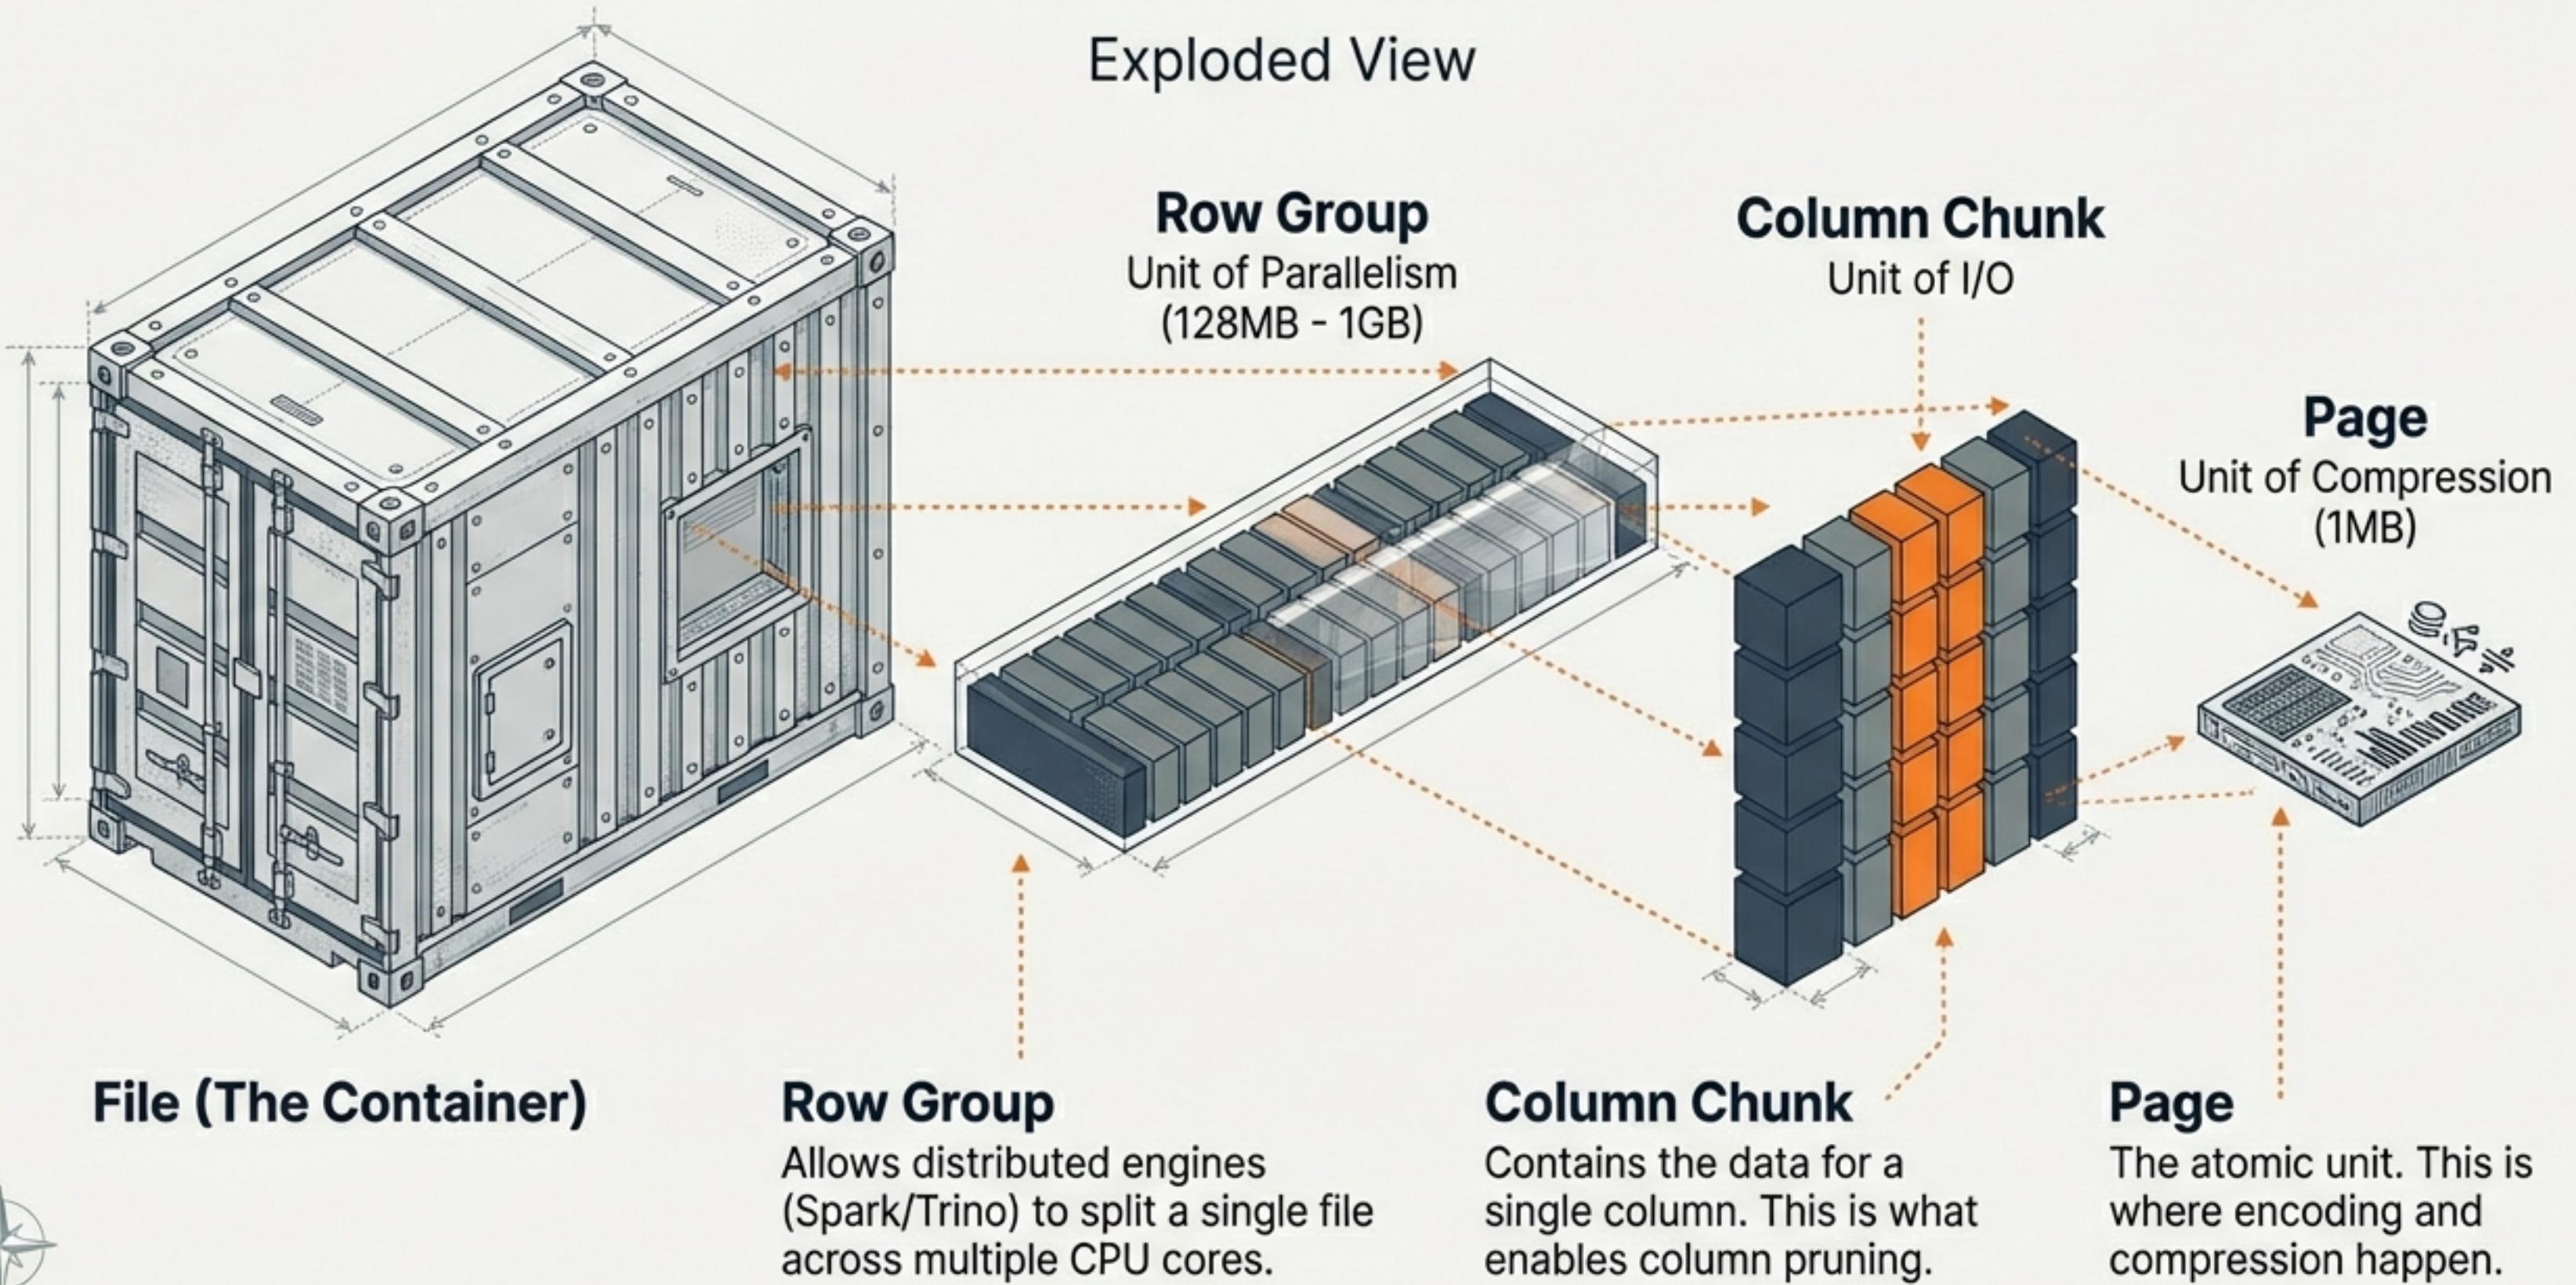


**99% I/O Reduction**

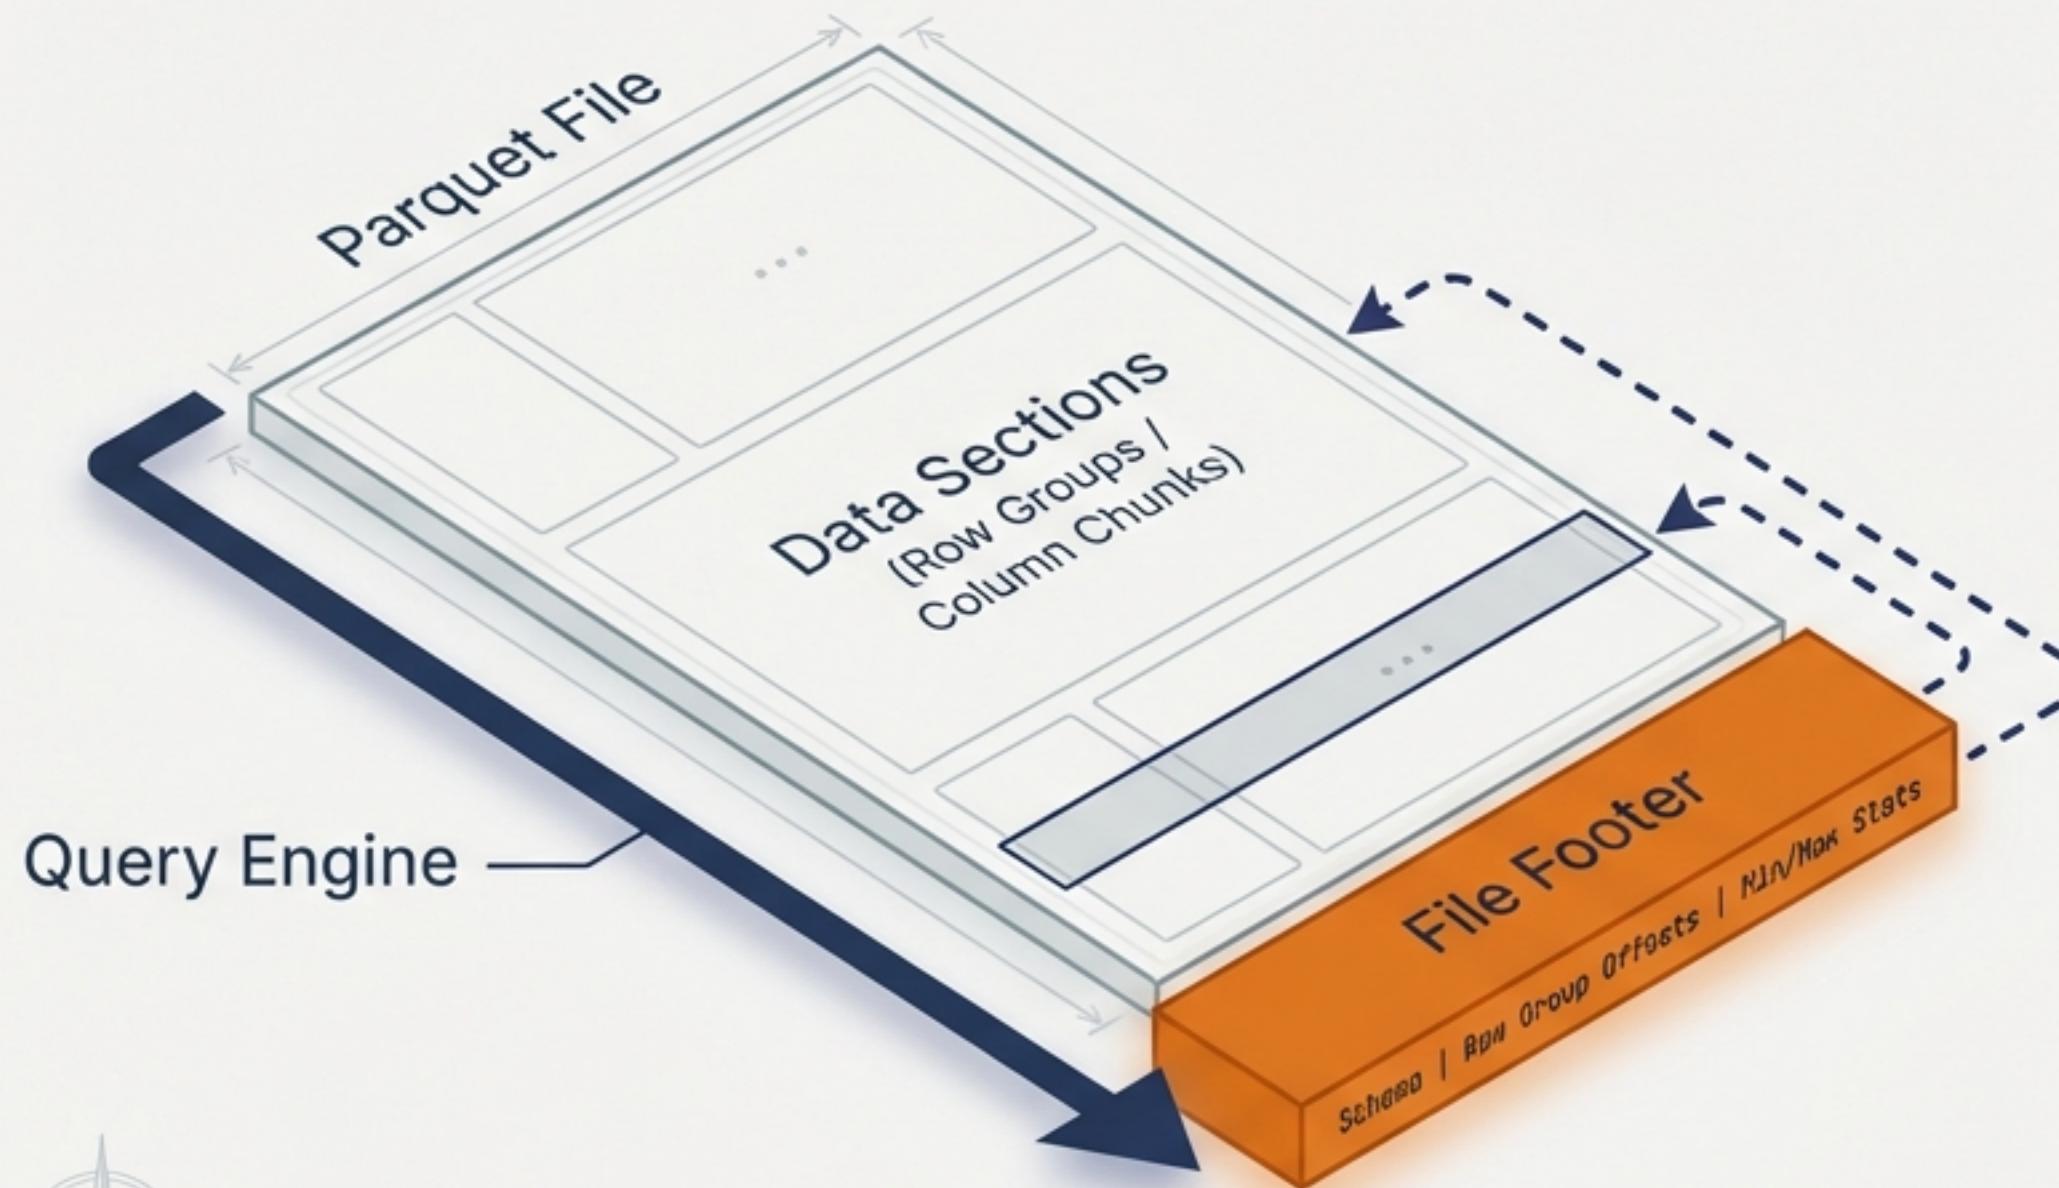
**Column Pruning:** In wide tables with hundreds of columns, analytic queries often touch only 3-4 metrics. Parquet physically skips the storage blocks of unrequested columns, meaning queries process gigabytes instead of terabytes.

# The Blueprint: Parquet File Hierarchy

Exploded View



# Architecture: Metadata First



## Smart Skipping

- Unlike CSV, Parquet is **self-describing**. The Footer contains the map. The query engine reads this map before processing any data.
- **Strategy**: If the metadata says a Row Group typically doesn't contain the requested data range, the engine never touches those bytes.

# The Art of Compression



**The Physics of Homogeneity:** Storing similar data types together maximizes the efficiency of compression codecs.

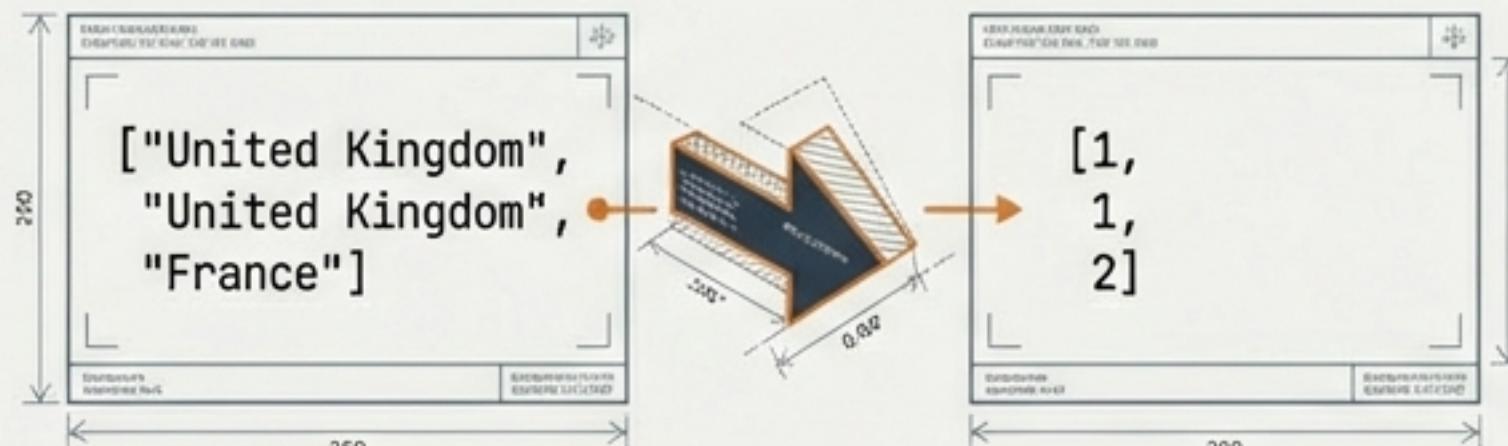
Codec Selection:

- **Snappy:** Optimised for speed (Hot Data).
- **Gzip/ZSTD:** Optimised for storage ratio (Cold Data).
- **ZSTD:** The modern gold standard for balancing both.

# Encoding: The Mathematics of Reduction

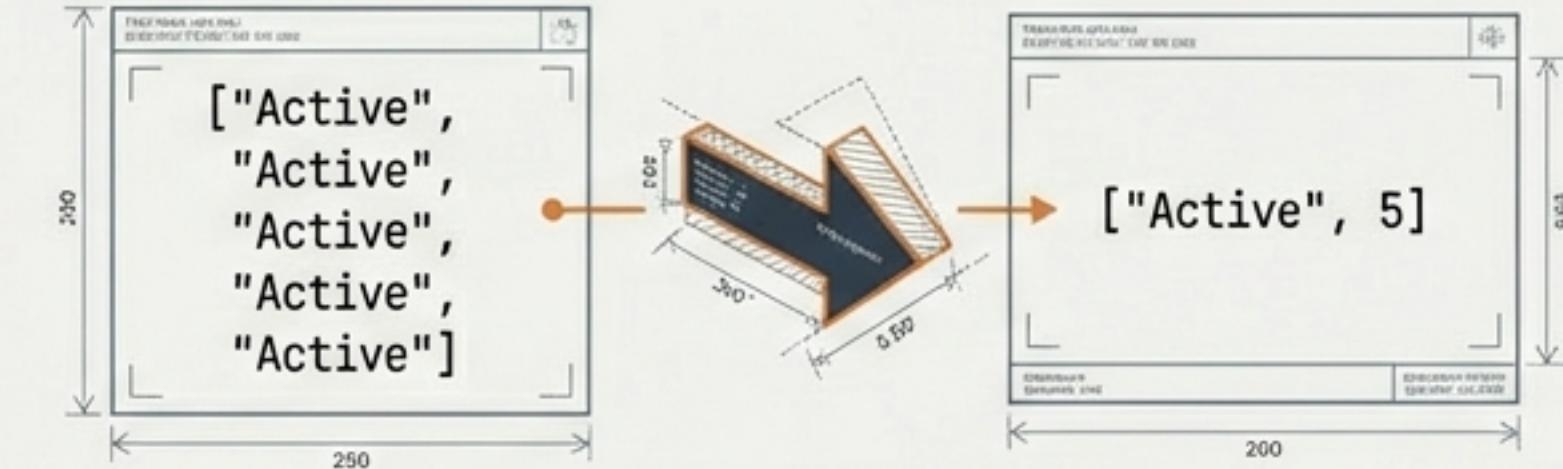
Logical reduction before physical compression.

## Dictionary Encoding



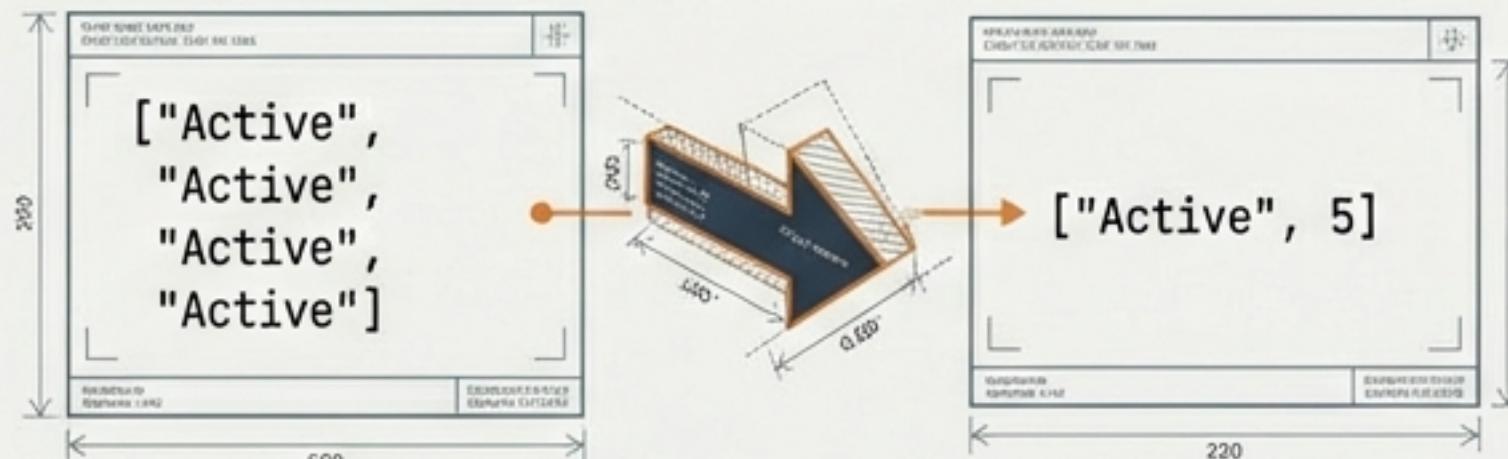
Replaces long strings with tiny integers.

## Run-Length Encoding (RLE)



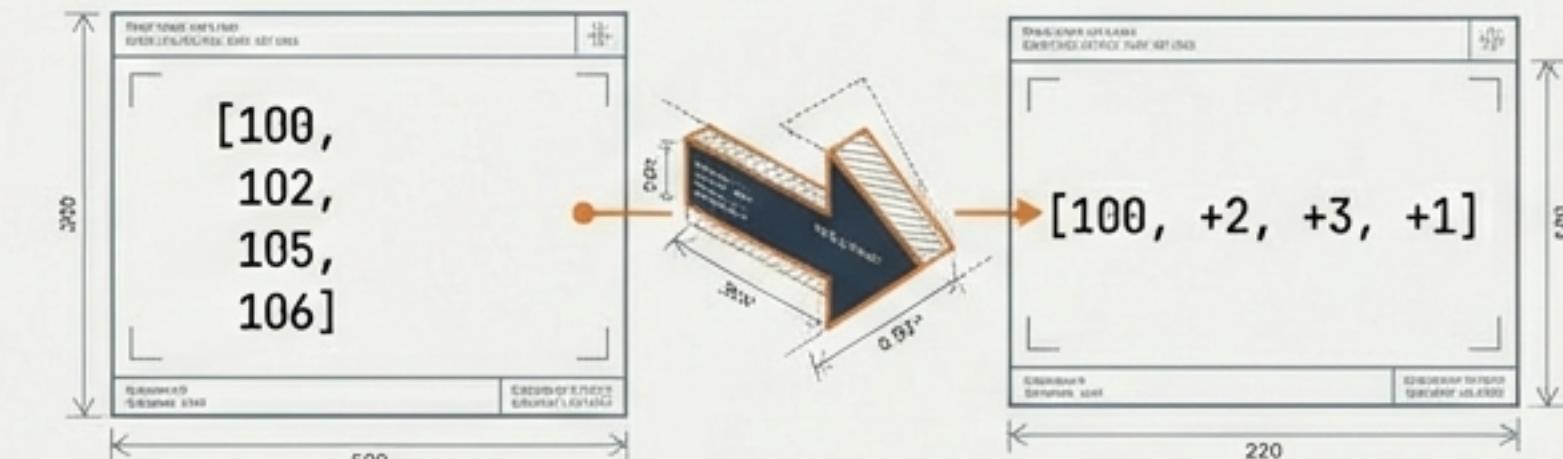
Collapses repetitive sequences.

## Eudon Encoding



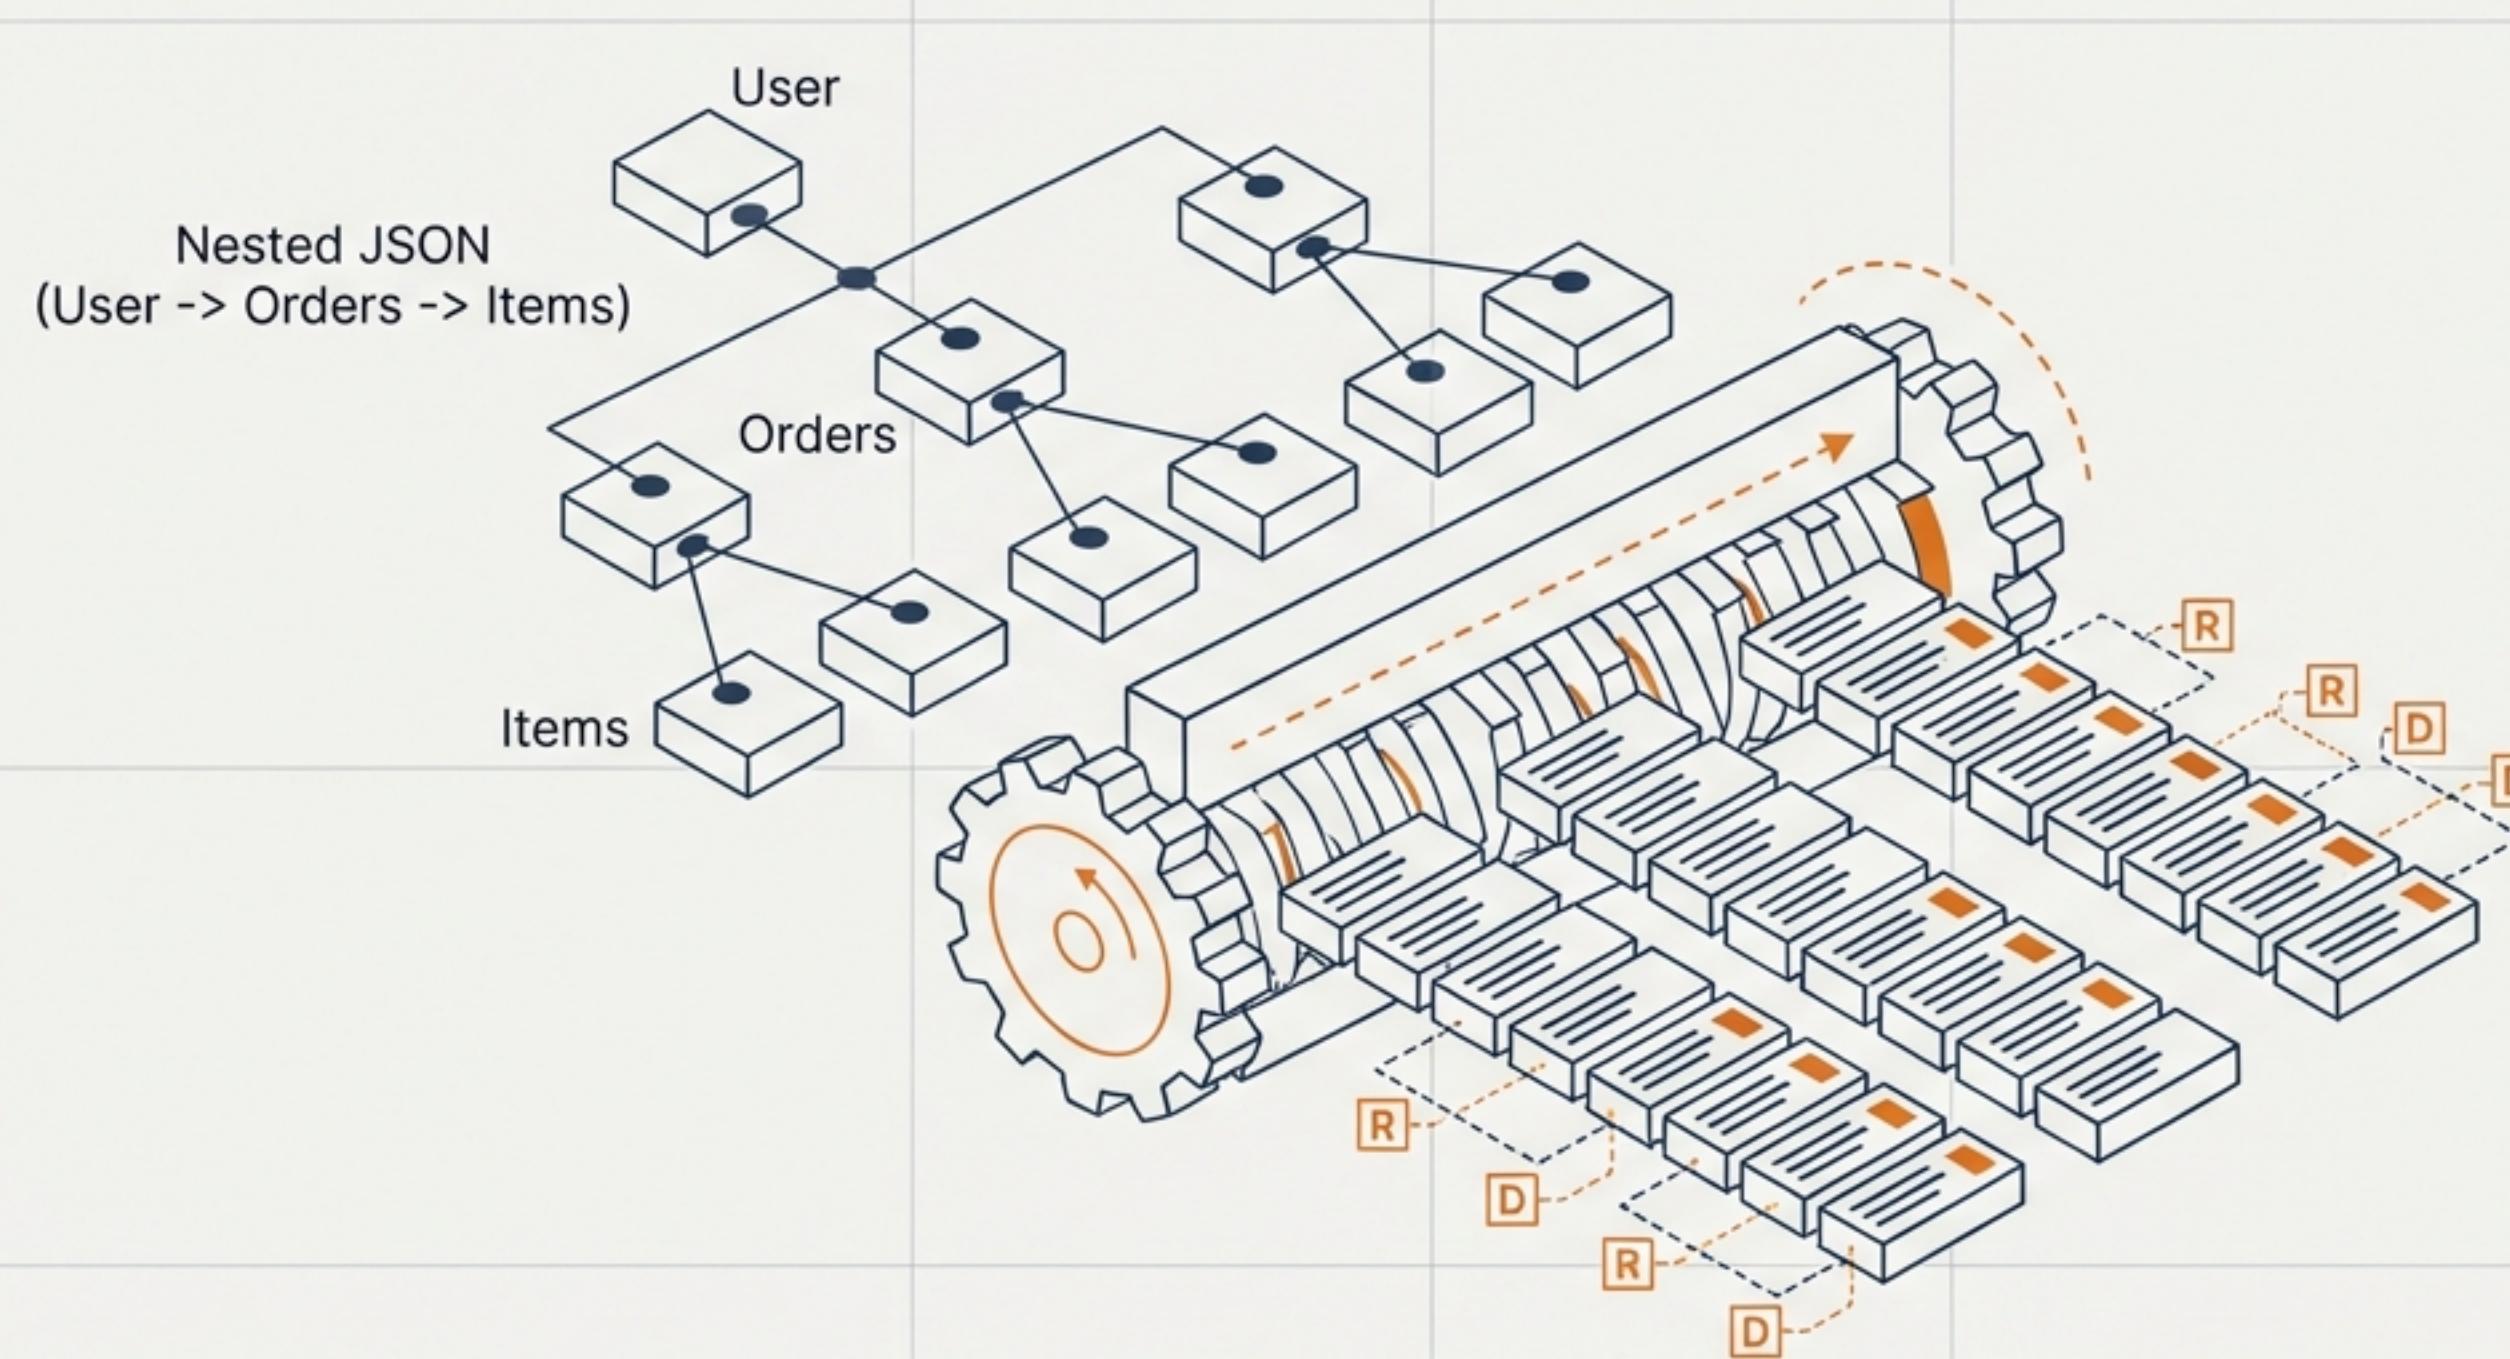
Collapses repetitive sequences.

## Delta Encoding



Stores small differences rather than large absolute values.

# Handling Complexity: The Dremel Algorithm



Precision  
Engineering  
Blueprint

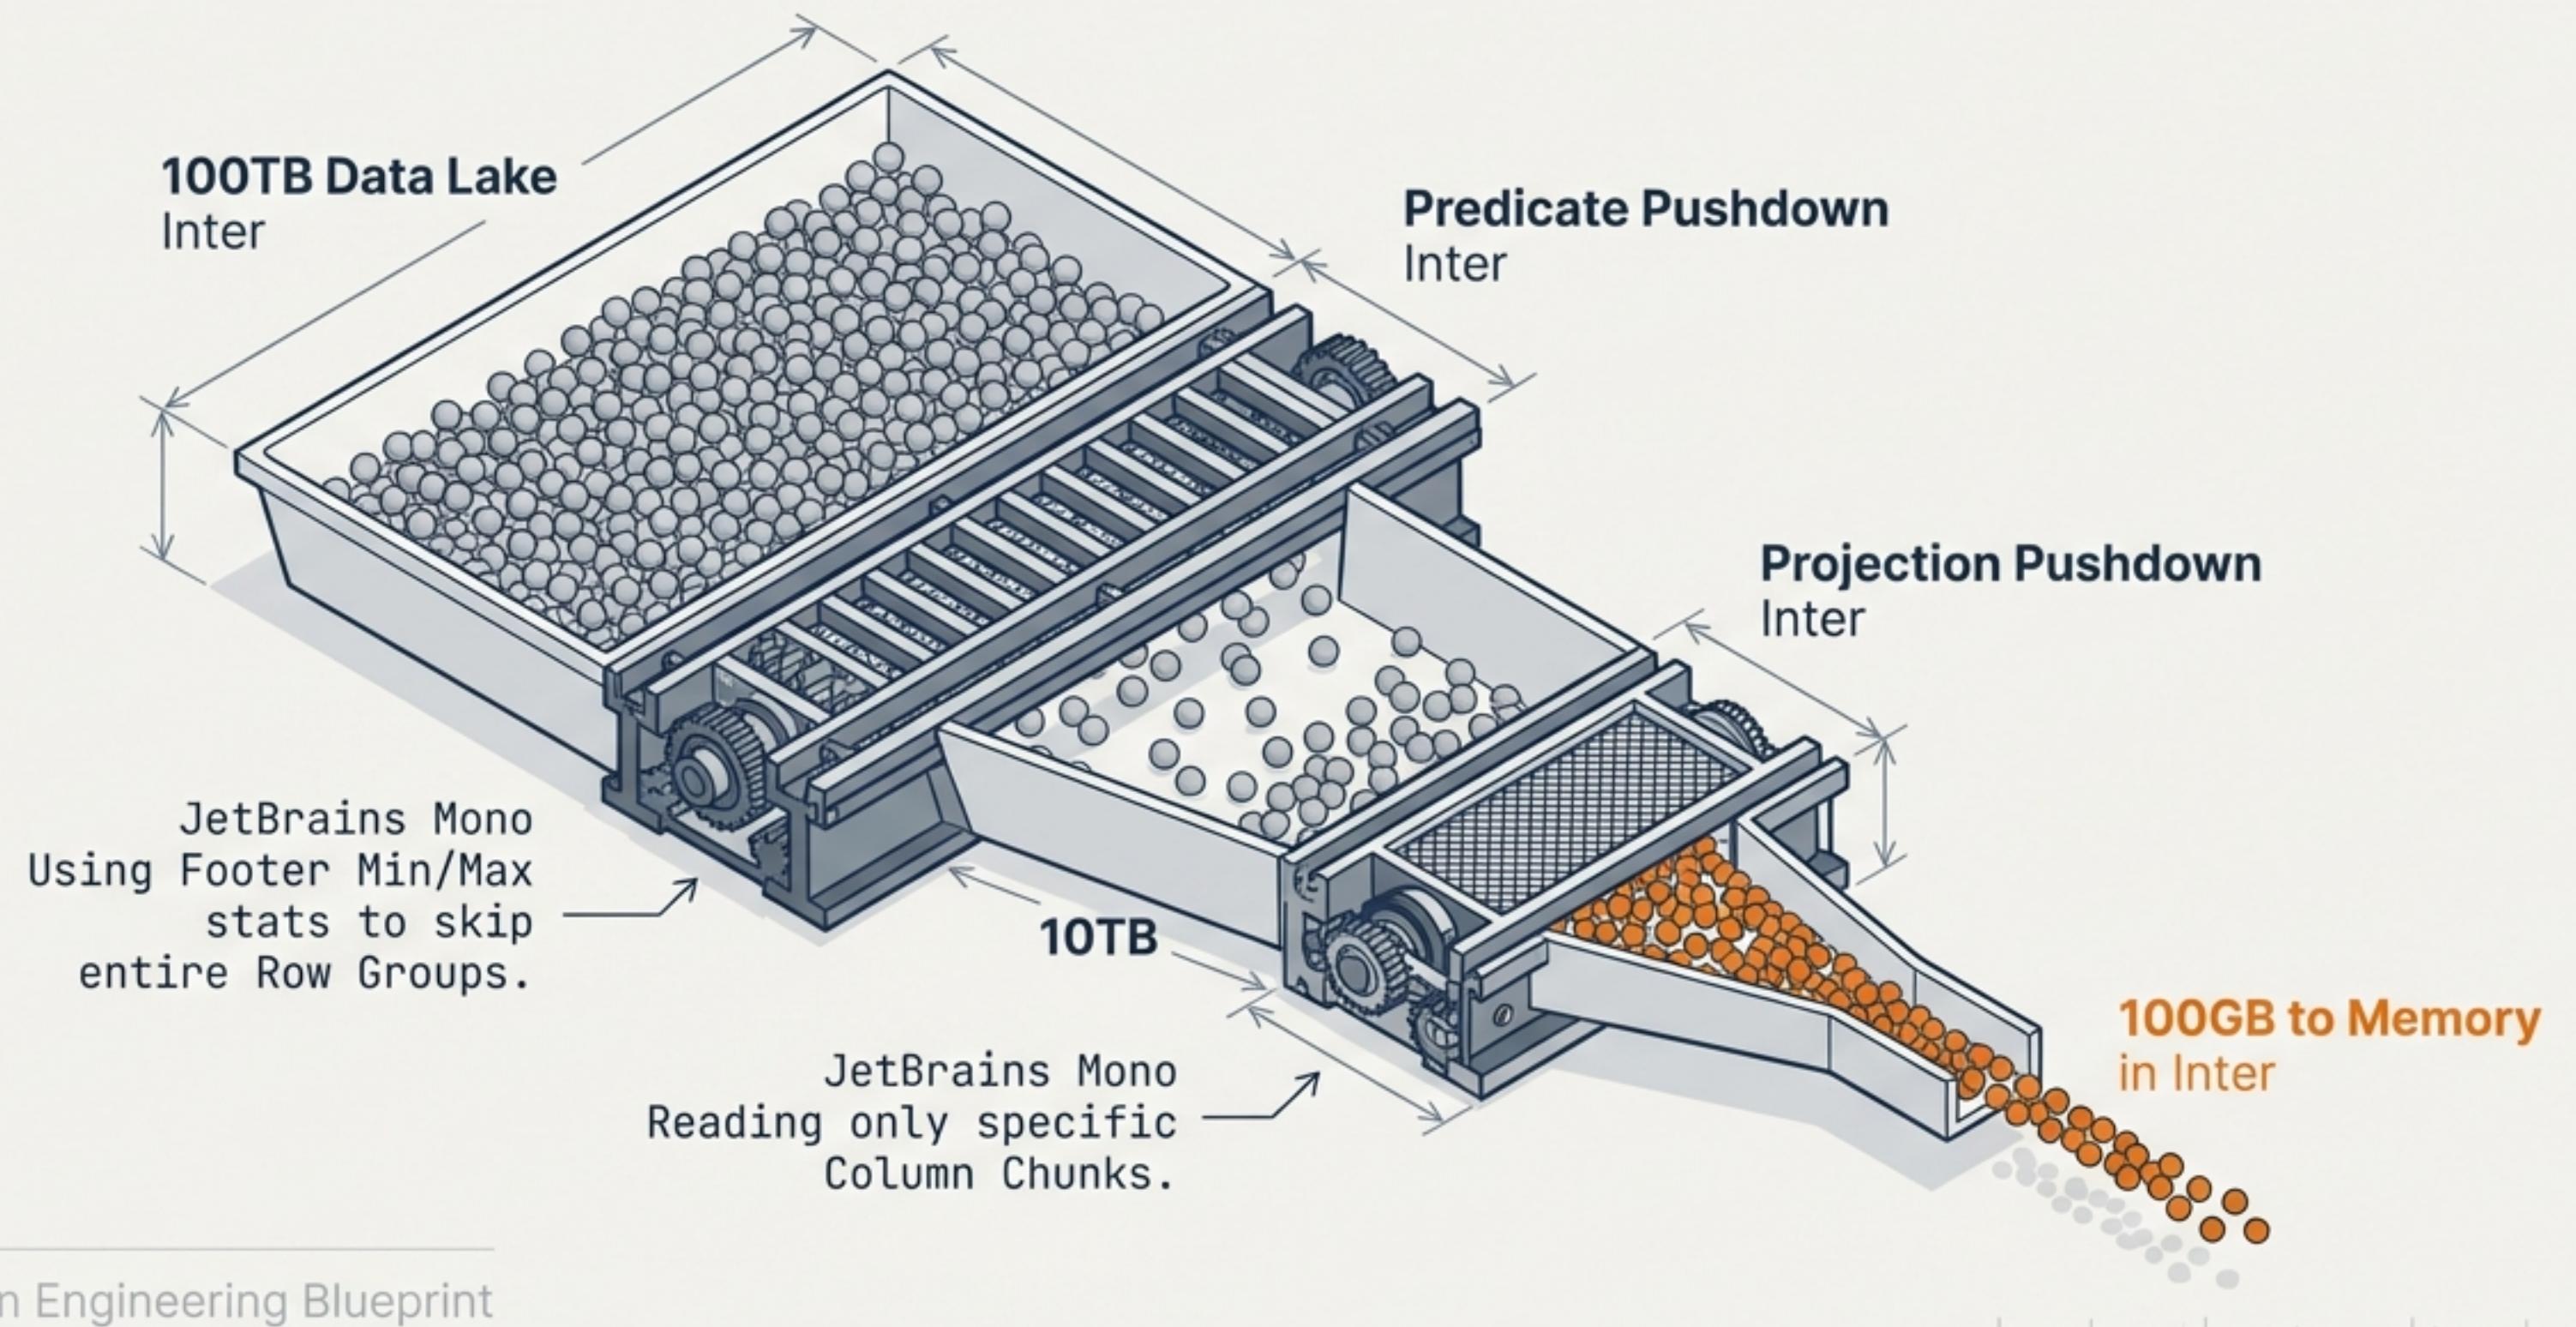
## Record Shredding

**Parquet** flattens nested structures (arrays/maps) into columns without data duplication.

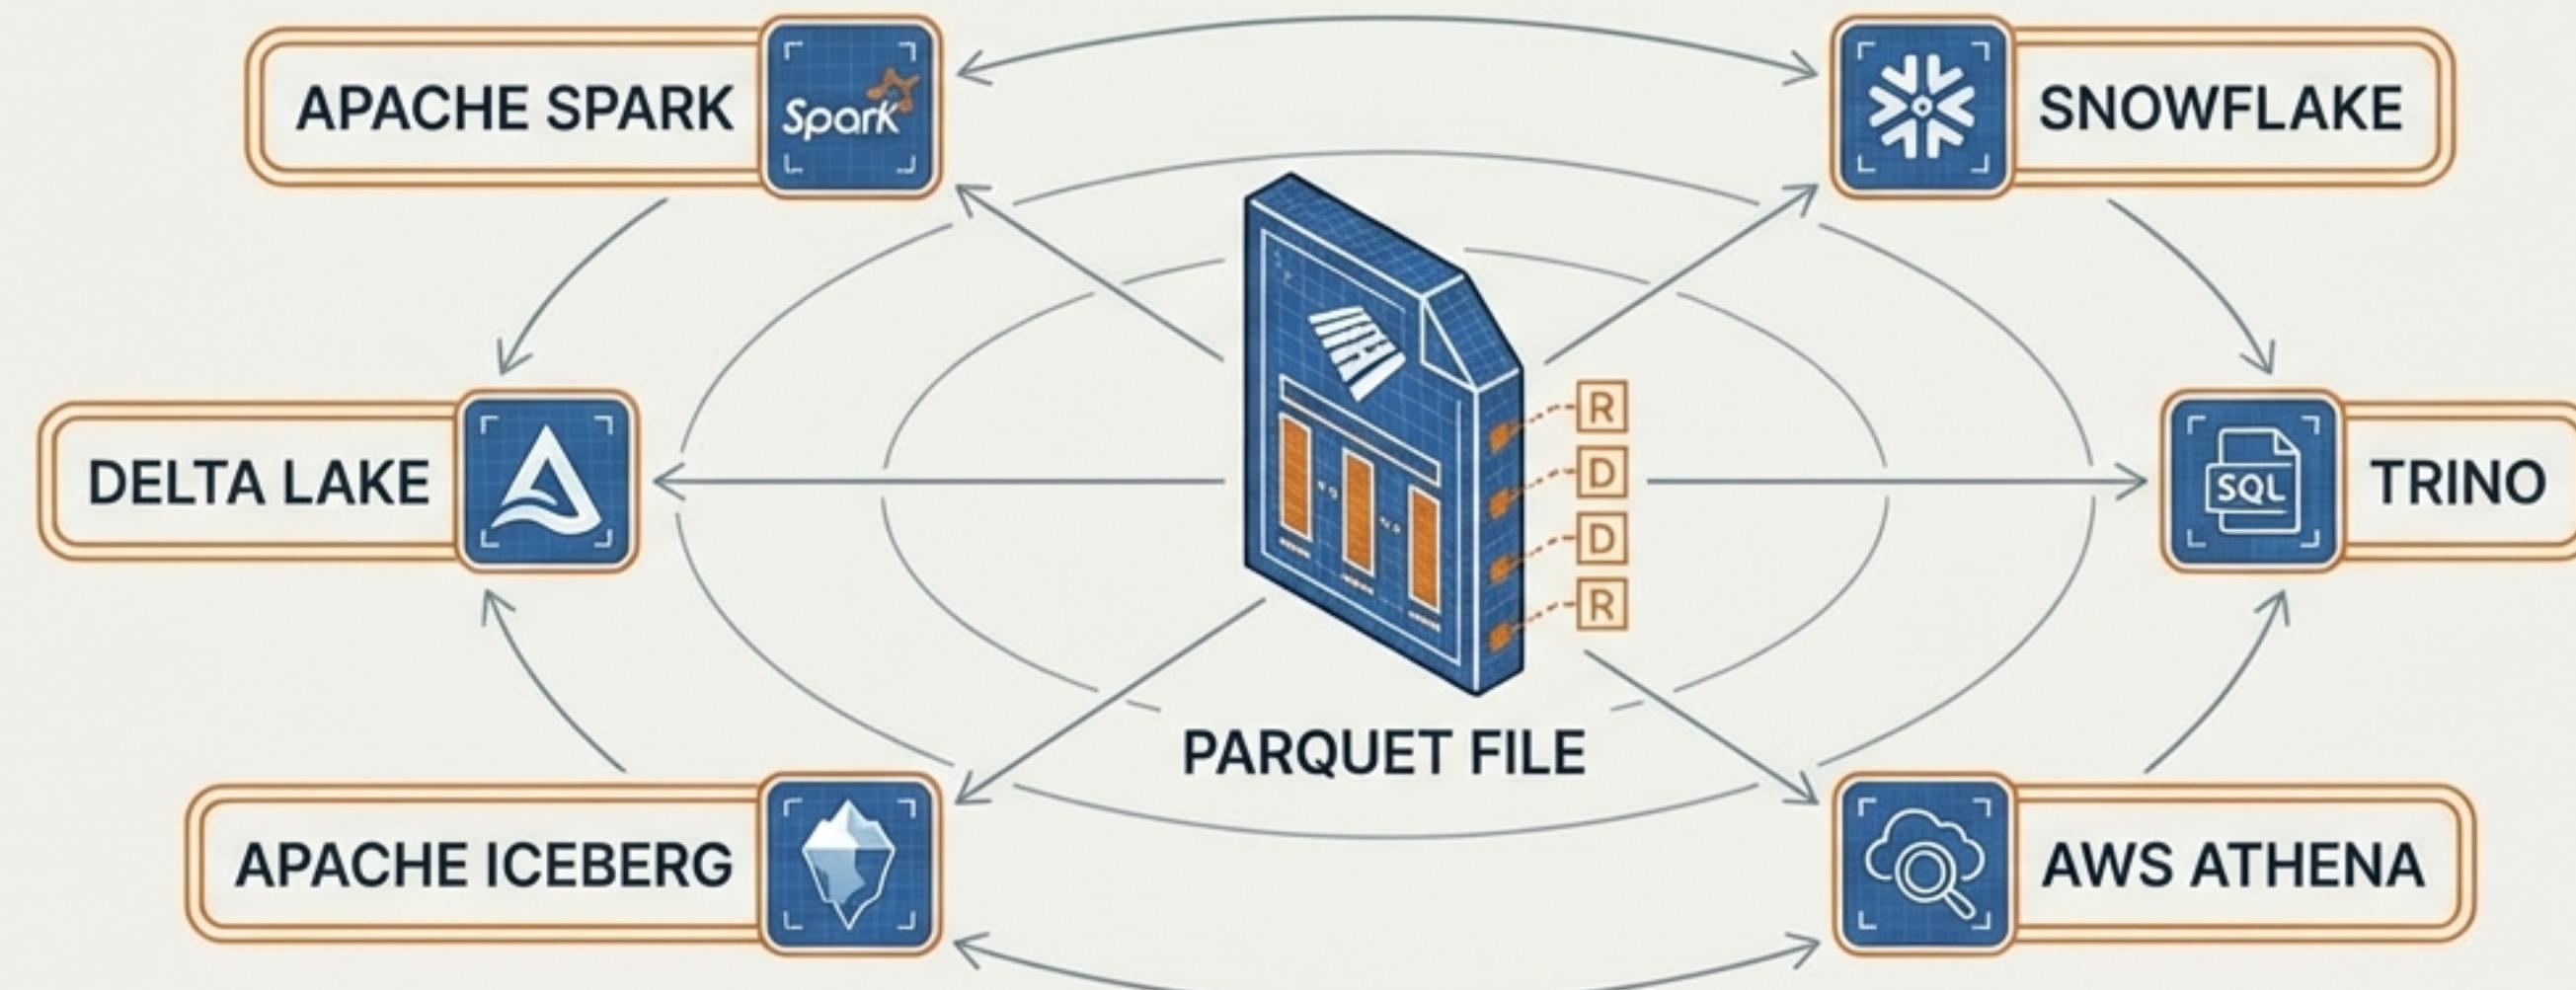
**Metadata Tags (R & D Levels):** Track depth and repetition, allowing the engine to reconstruct the original nested structure on read.

**Benefit:** Query a single field deep inside a complex object without parsing the entire tree.

# Optimization: Don't Pay for Data You Don't Use



# The Universal Currency of the Lakehouse

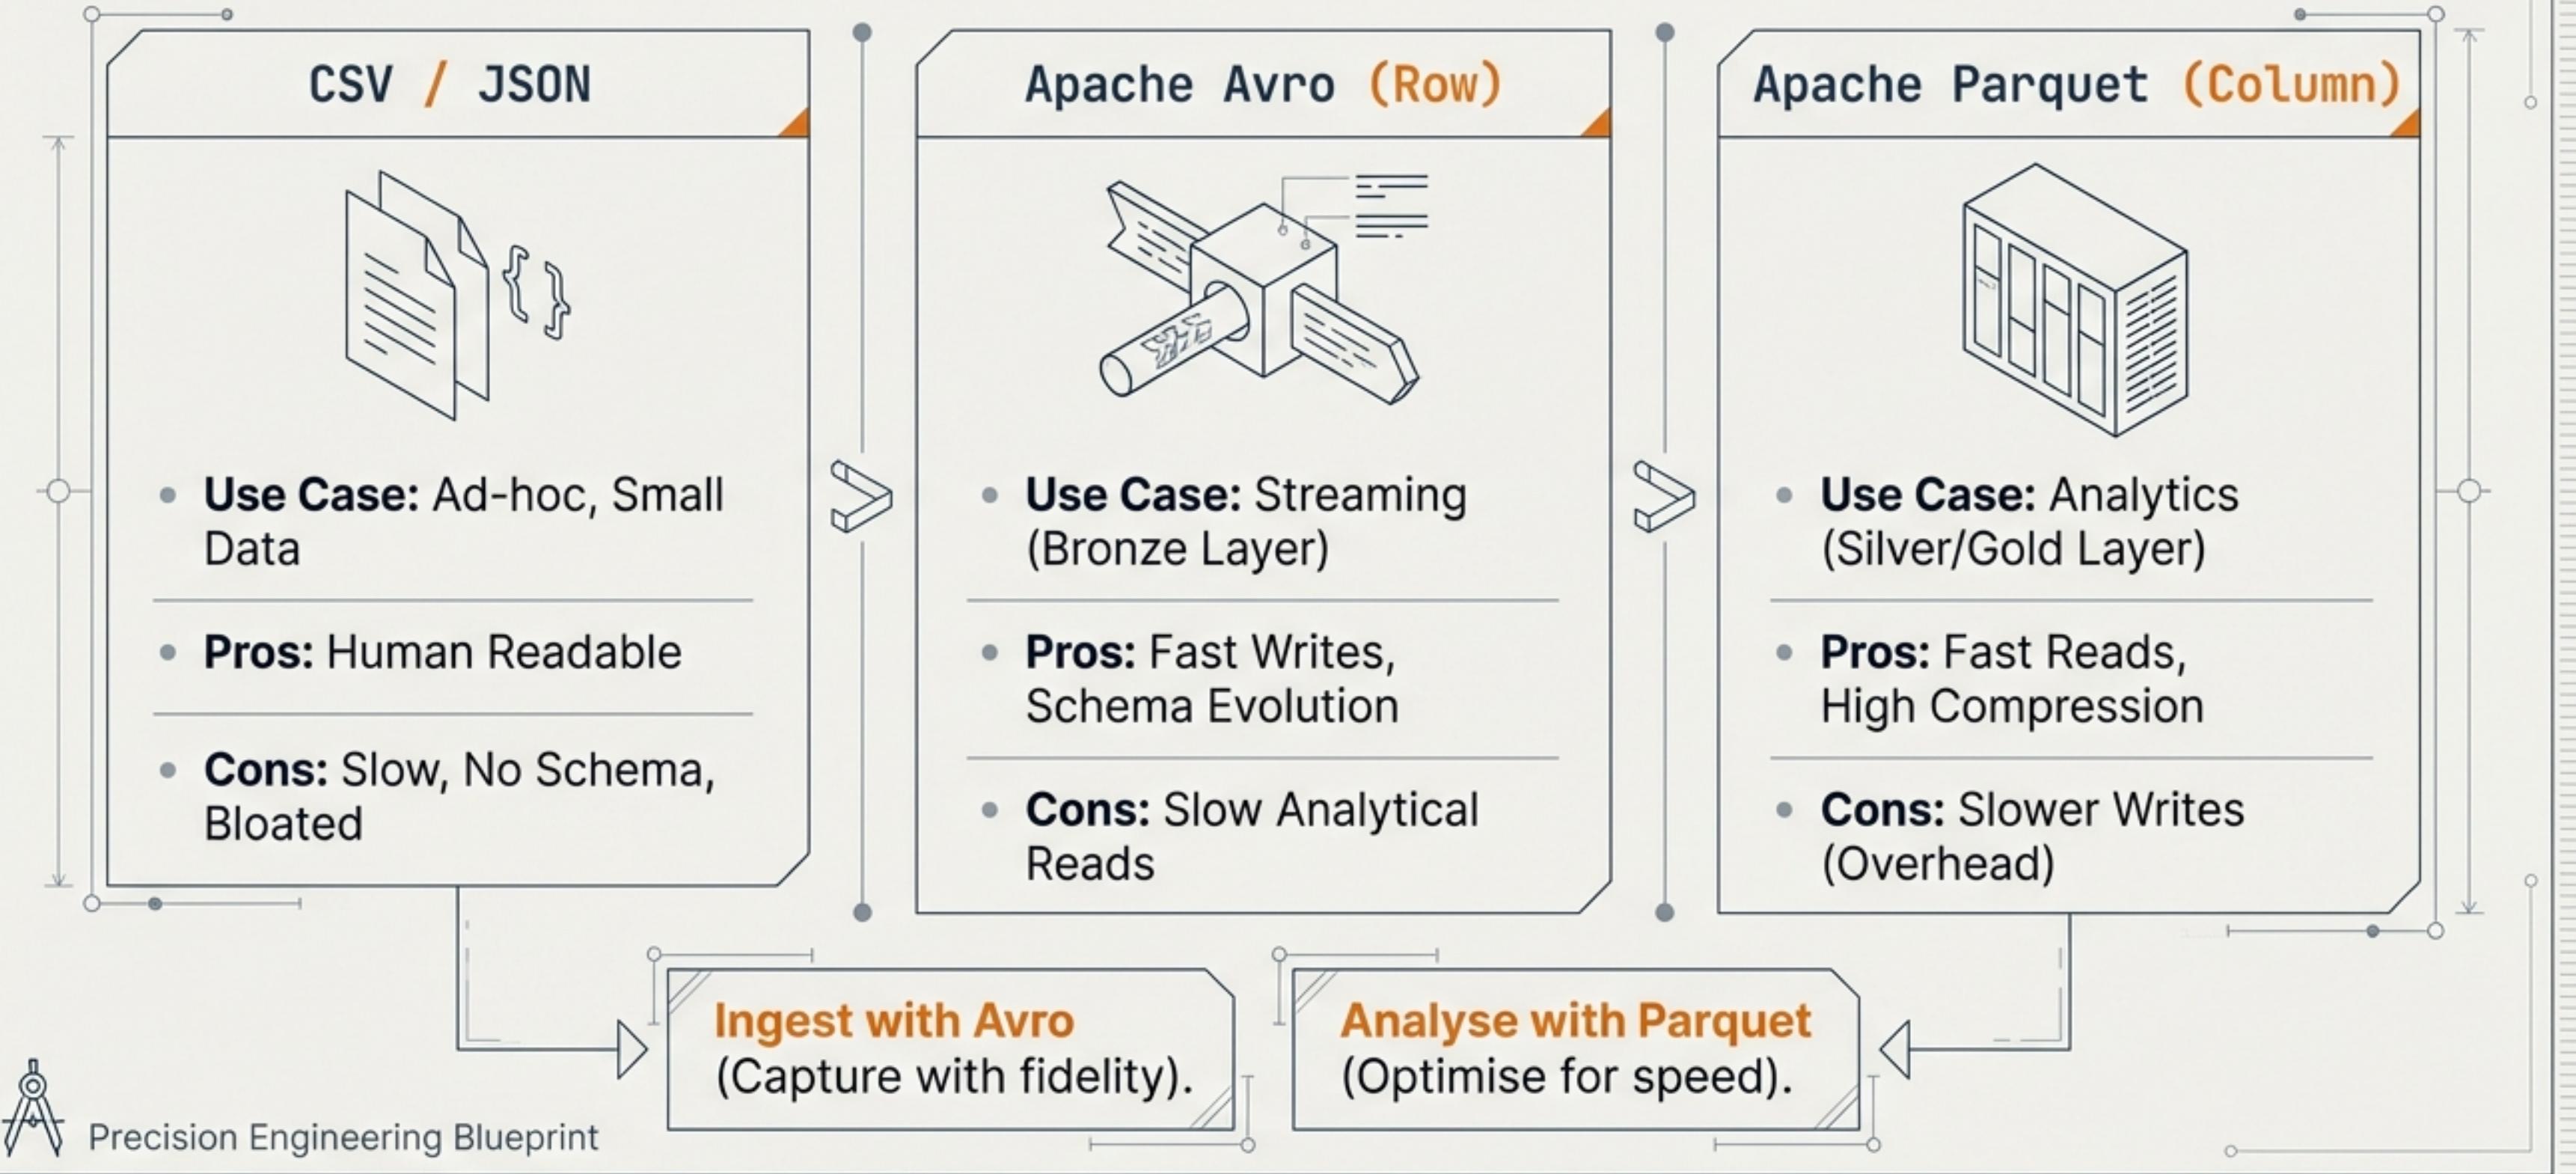


**Decoupled Architecture:** Write data once with **Spark**, read it later with **Snowflake** or **Trino**. No conversion required.

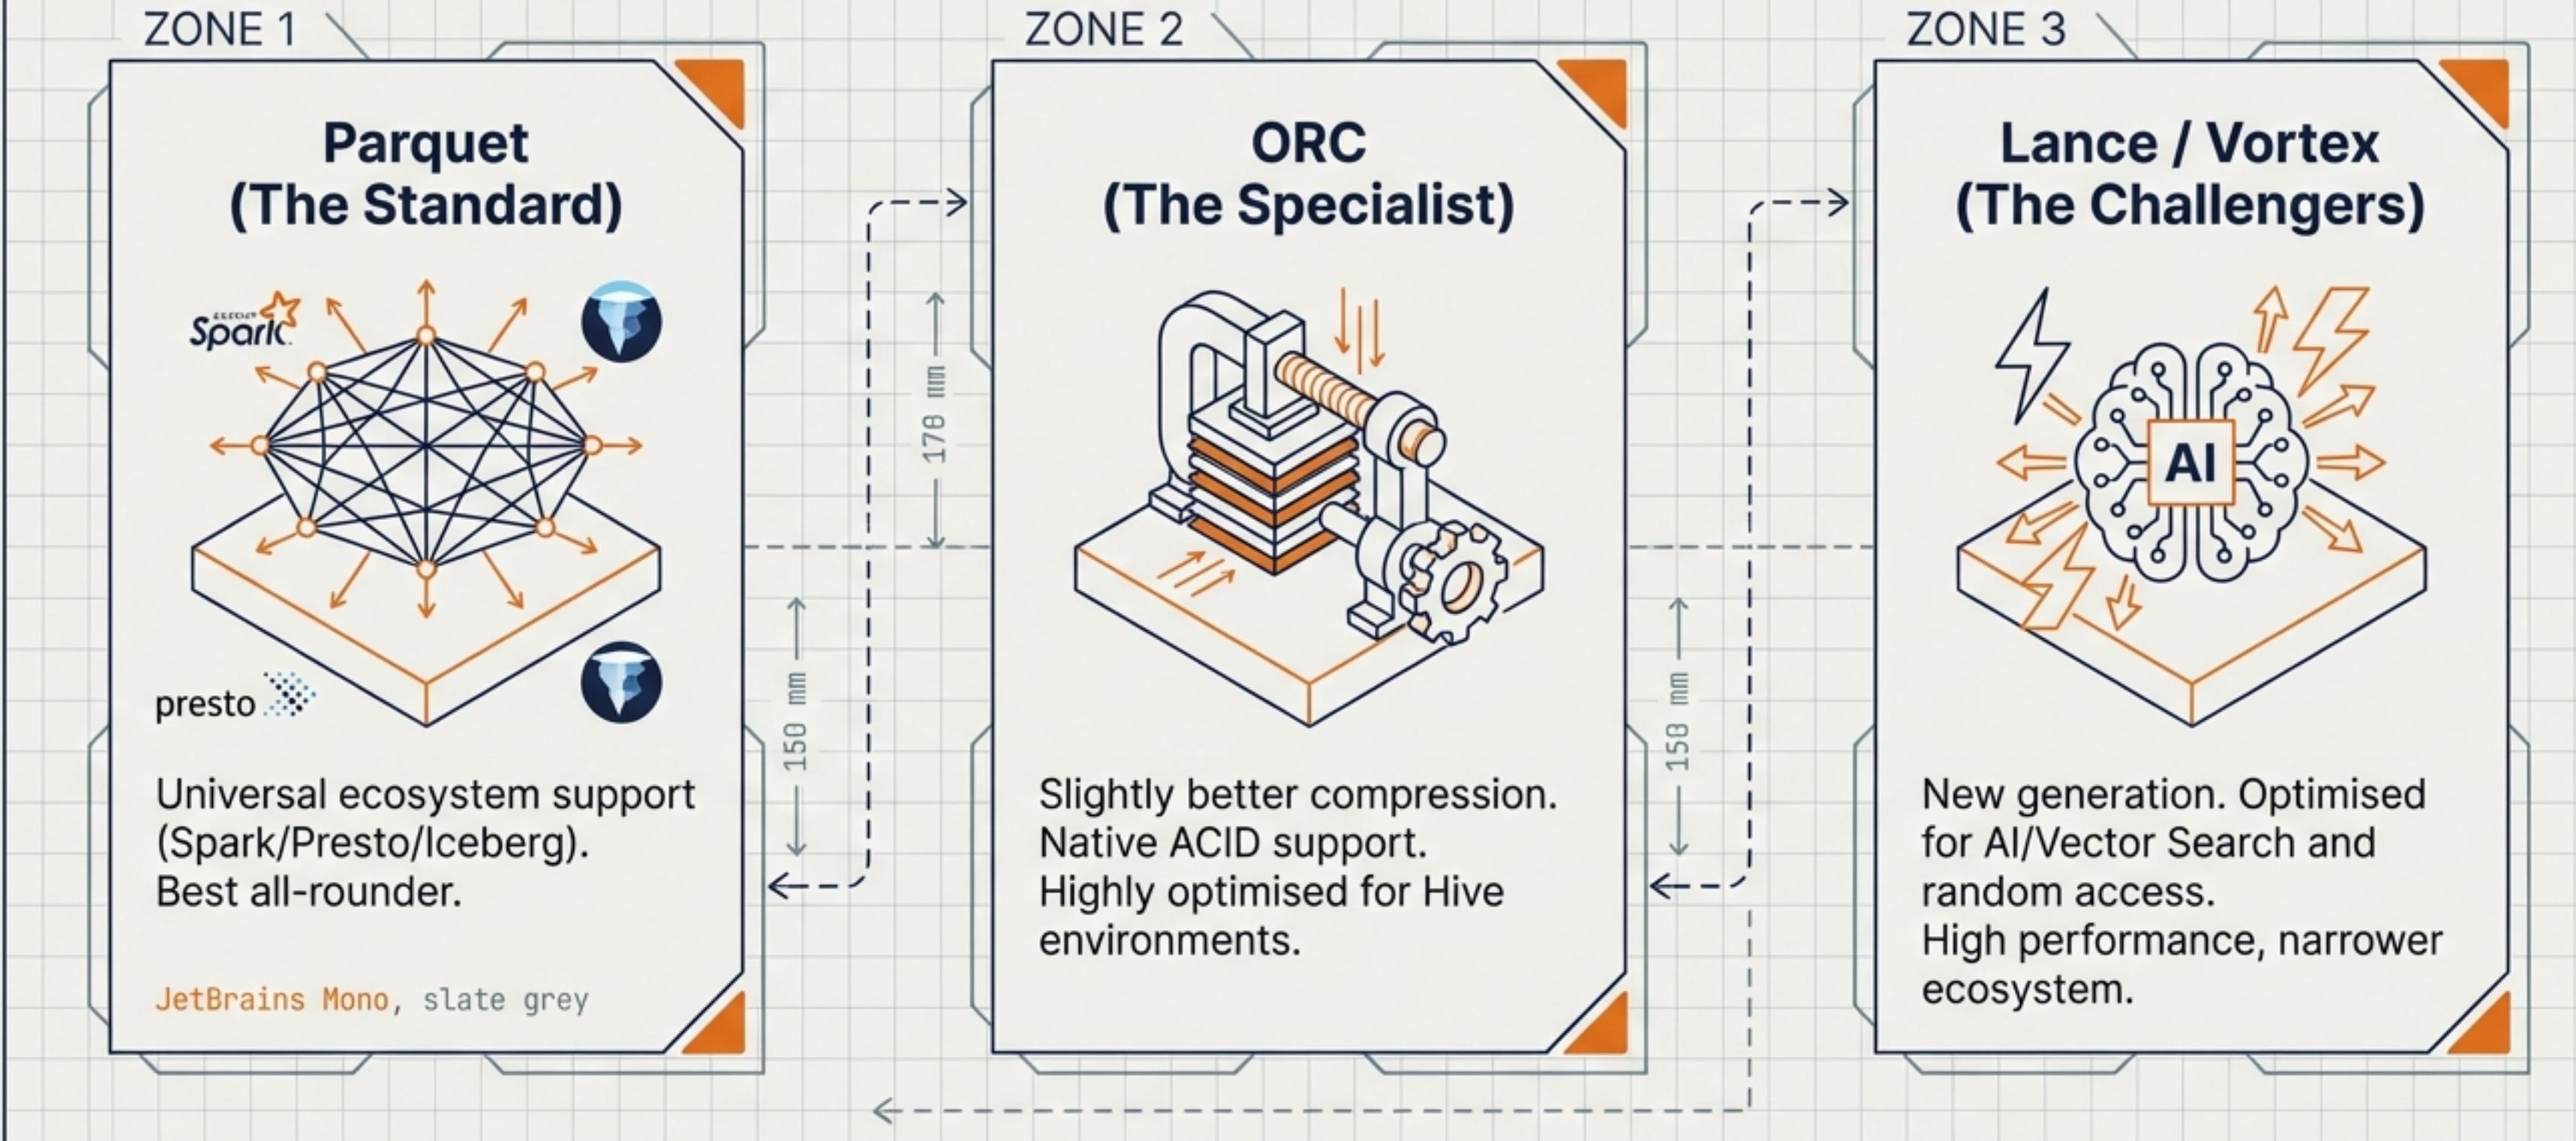
**Interoperability:** The foundational storage layer for modern table formats like **Iceberg** and **Delta Lake**.



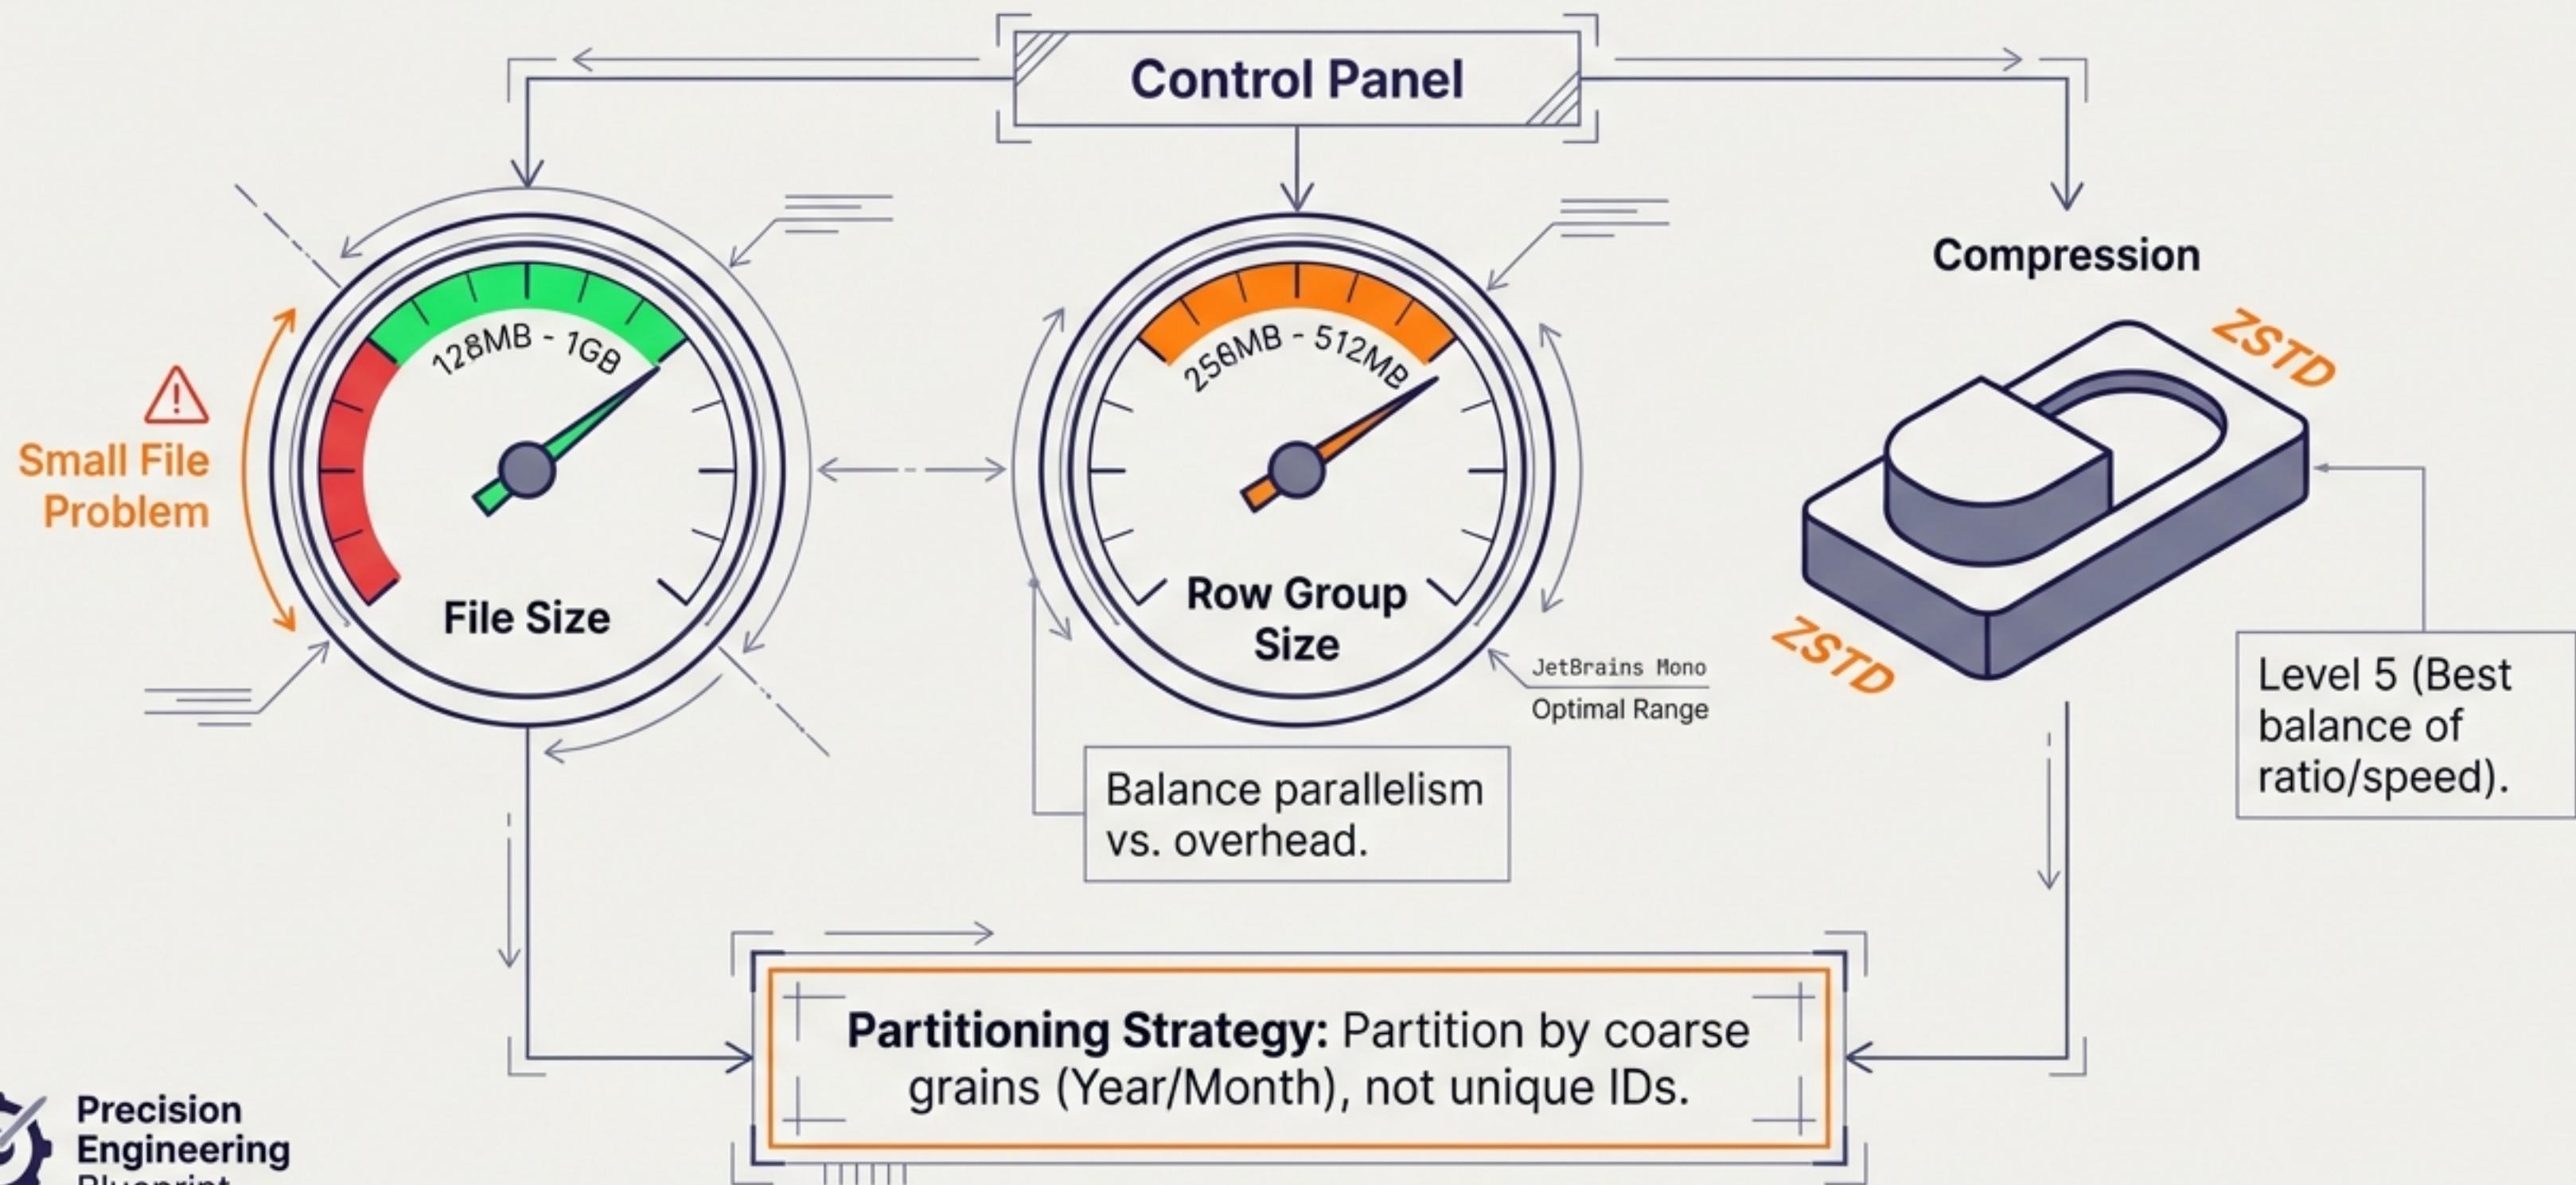
# The Showdown: Row vs. Columnar



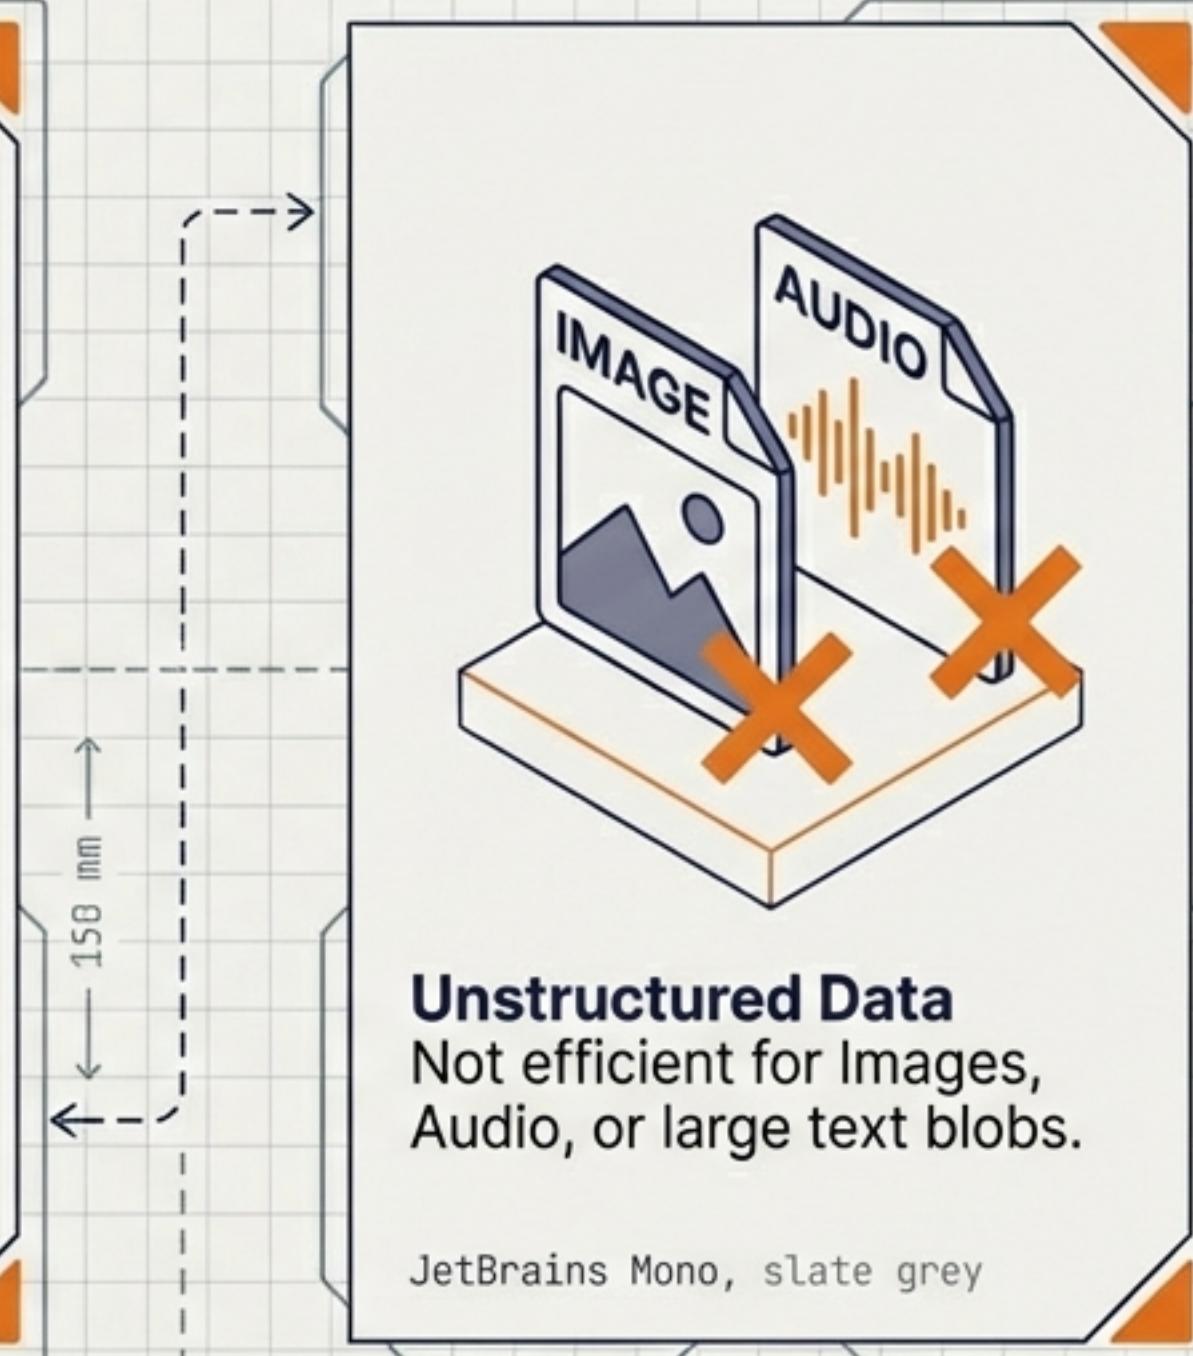
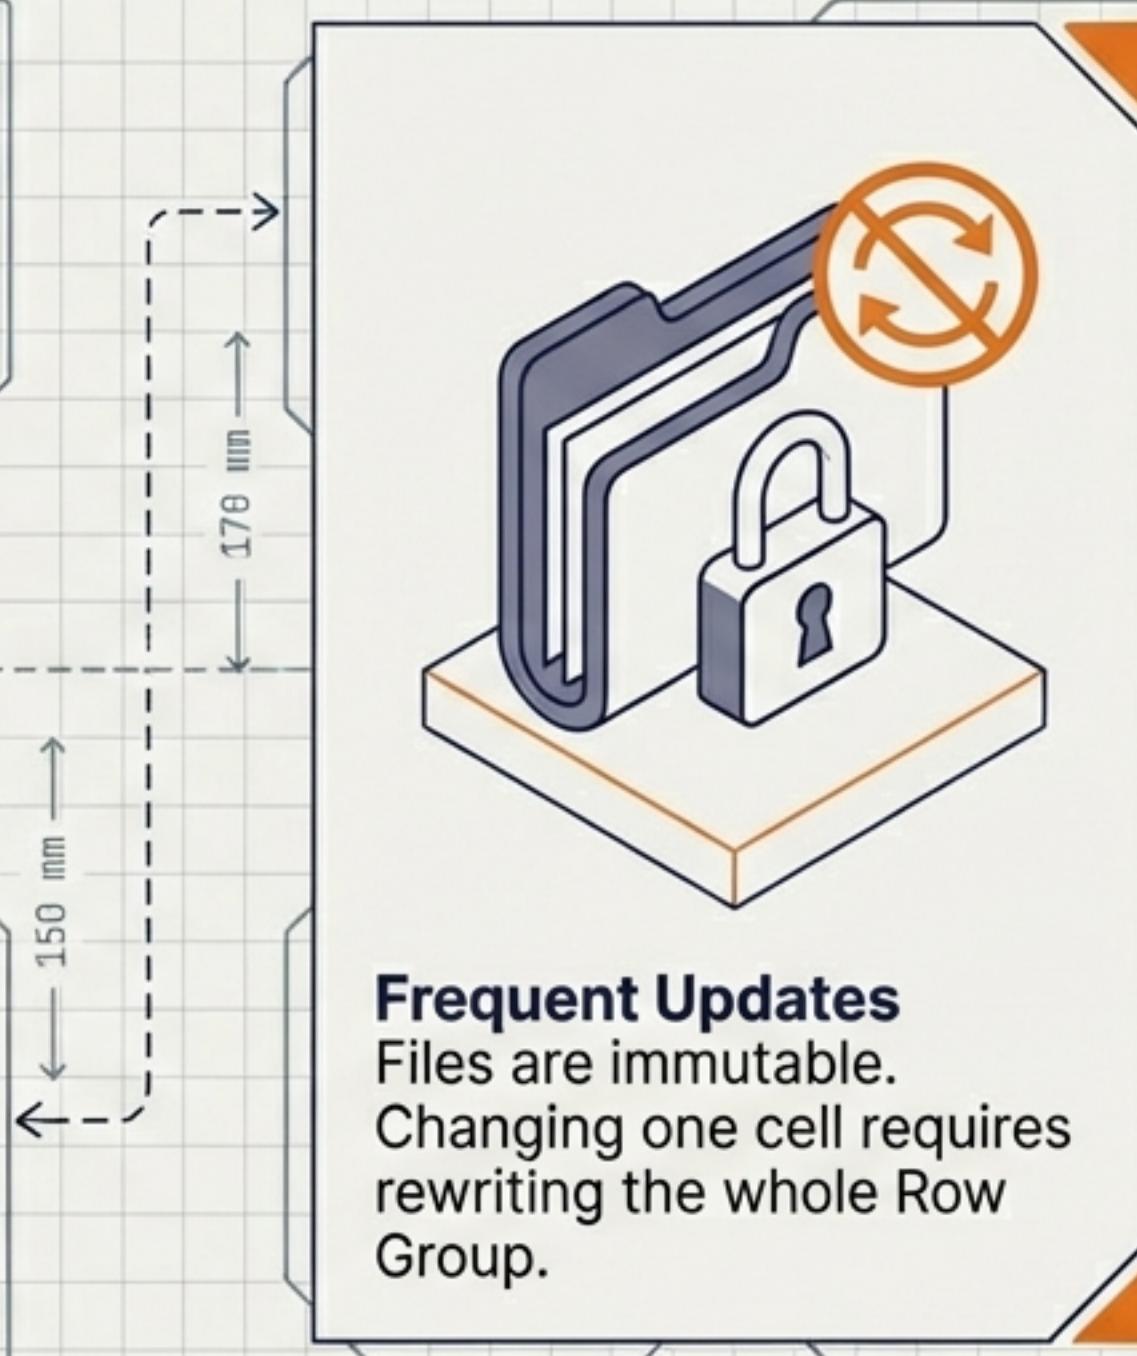
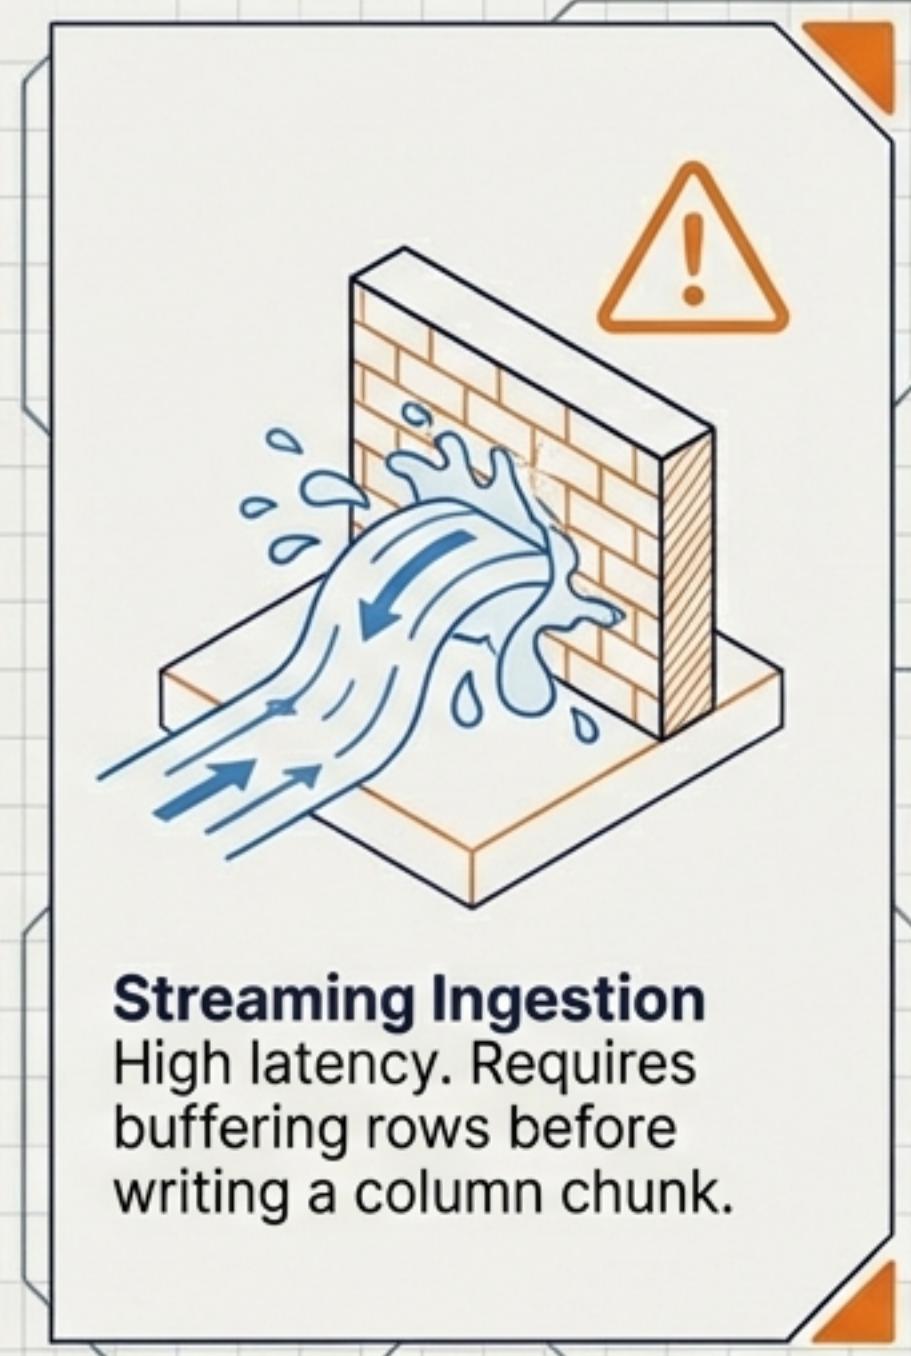
# The Columnar Landscape



# Best Practices: Tuning for Performance



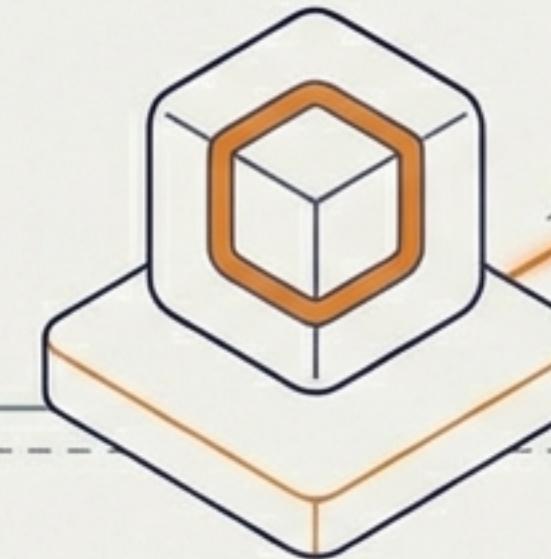
# When NOT to Use Parquet



JetBrains Mono, slate grey

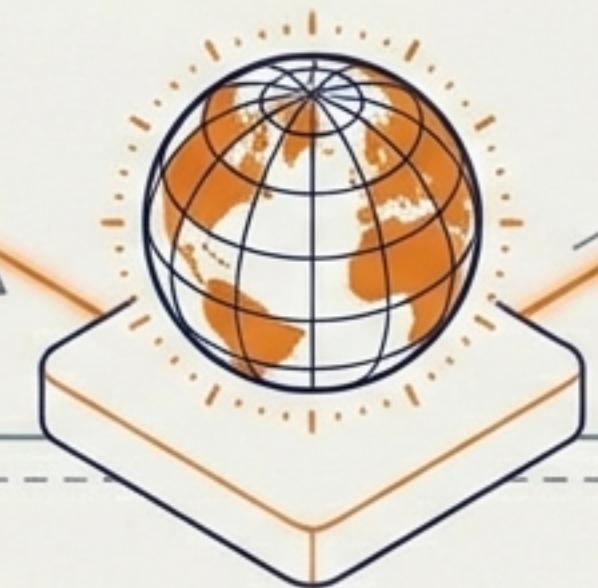
# The Future of Parquet

**Variant Type**



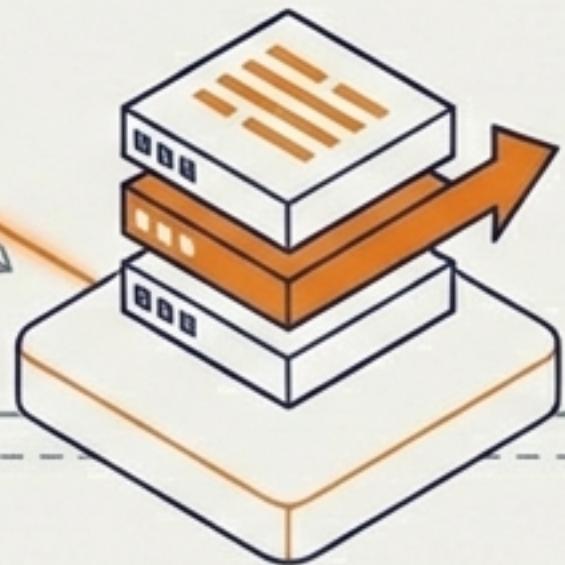
Native support for  
semi-structured JSON without  
string parsing overhead.

**GeoParquet**



Standardised encoding  
for geospatial data.

**FlatBuffers**



Accelerated metadata  
parsing.

**The bedrock of modern  
data engineering.**

